

Docker

Basic Docker Components

Courtesy John Willis



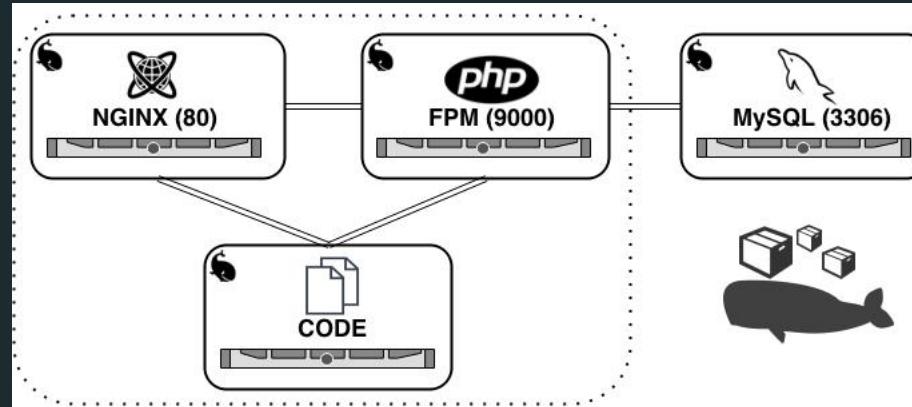
Today

- Introduction to containers
- Installing Docker
- Docker concepts and terms
- Introduction to images
- Running and managing containers
- Building images
- Container volumes
- Container networking



Workshop

- Docker Setup
- Docker RUN
- Docker Compose



Module 1: Introduction to Containers



Module objectives

In this module we will:

- Introduce the concept of container based virtualization
- Outline the benefits of containers over virtual machines



What is Docker?

Docker is a platform for developing, shipping and running applications using container technology

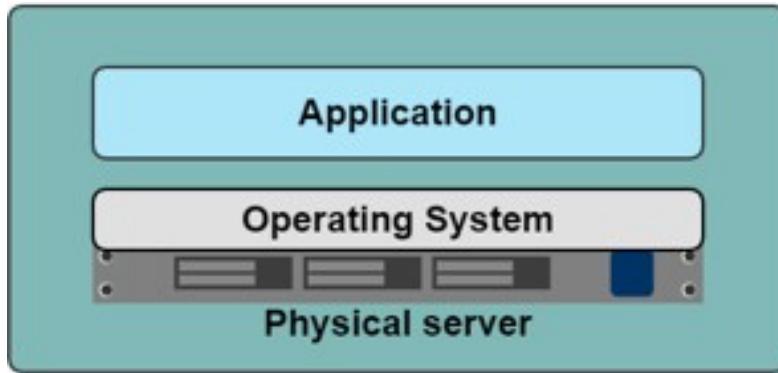
- The Docker Platform consists of multiple products/tools
 - Docker Engine
 - Docker Hub
 - Docker Trusted Registry
 - Docker Machine
 - Docker Swarm
 - Docker Compose



A History Lesson

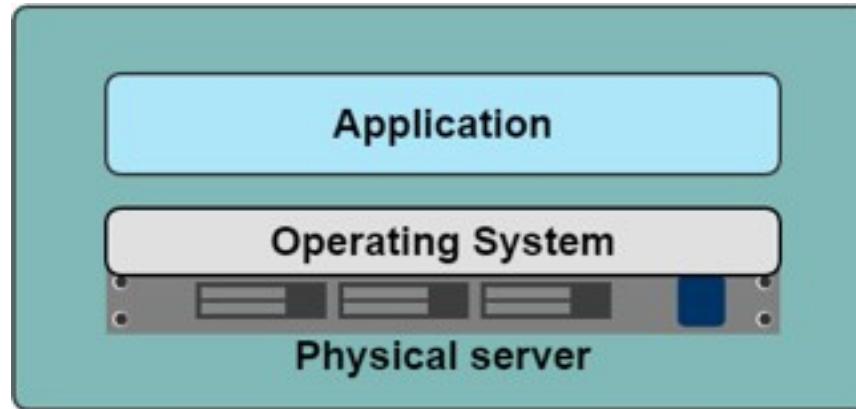
In the Dark Ages

One application on one physical server



Historical limitations of application deployment

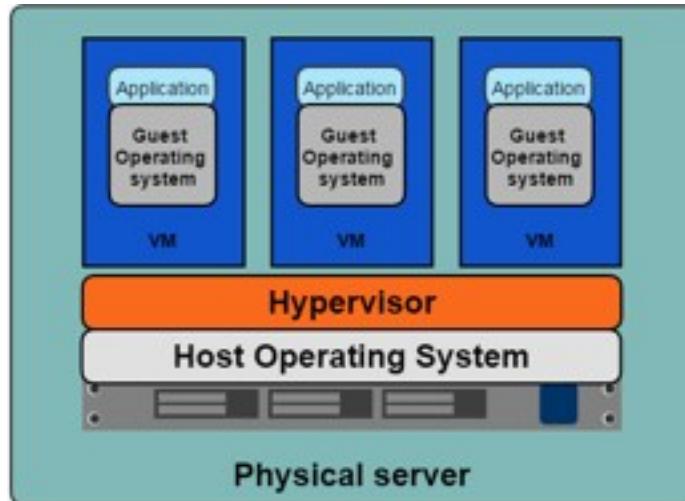
- Slow deployment times
- Huge costs
- Wasted resources
- Difficult to scale
- Difficult to migrate
- Vendor lock in



A History Lesson

Hypervisor-based Virtualization

- One physical server can contain multiple applications
- Each application runs in a virtual machine (VM)



Benefits of VM's

- Better resource pooling
 - One physical machine divided into multiple virtual machines
- Easier to scale
- VM's in the cloud
 - Rapid elasticity
 - Pay as you go model



Google Cloud



Limitations of VM's

- Each VM stills requires
 - CPU allocation
 - Storage
 - RAM
 - An entire guest operating system
- The more VM's you run, the more resources you need
- Guest OS means wasted resources
- Application portability not guaranteed



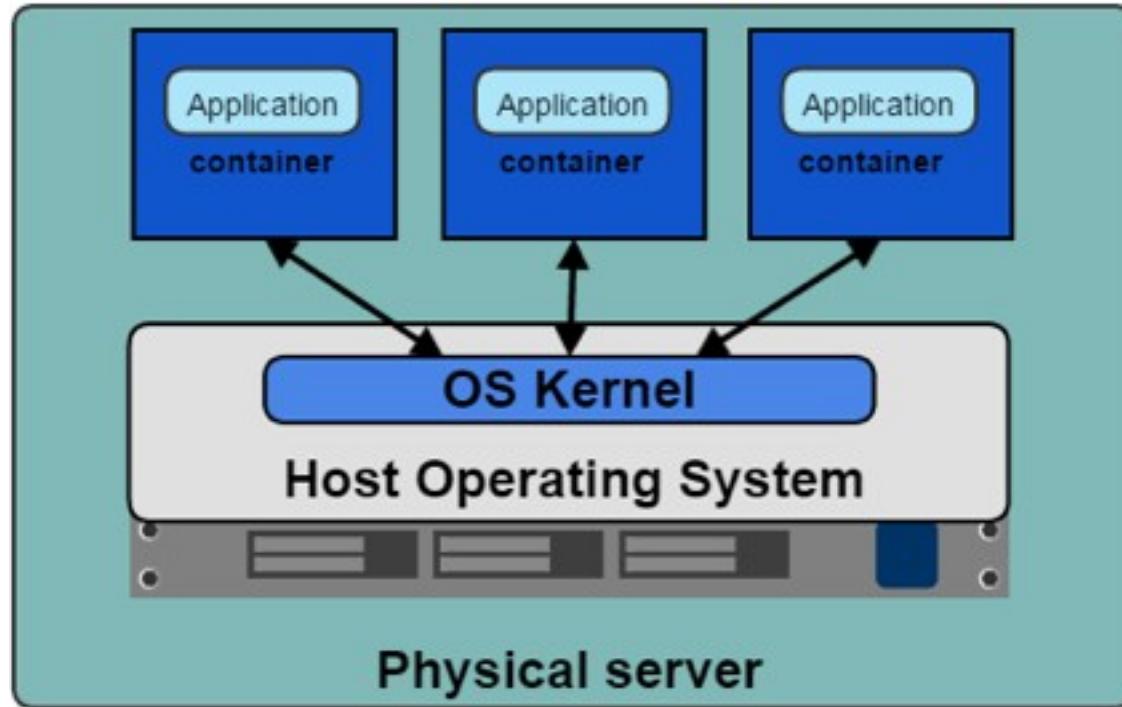
Introducing Containers

Containerization uses the kernel on the host operating system to run multiple root file systems

- Each root file system is called a **container**
- Each container also has its own
 - Processes
 - Memory
 - Devices
 - Network stack



Containers



Containers vs VM's

- Containers are more lightweight
- No need to install guest OS
- Less CPU, RAM, storage space required
- More containers per machine than VMs
- Speed of instantiation
- Greater portability



Why Docker?

- Isolation
- Lightweight
- Simplicity
- Workflow
- Community



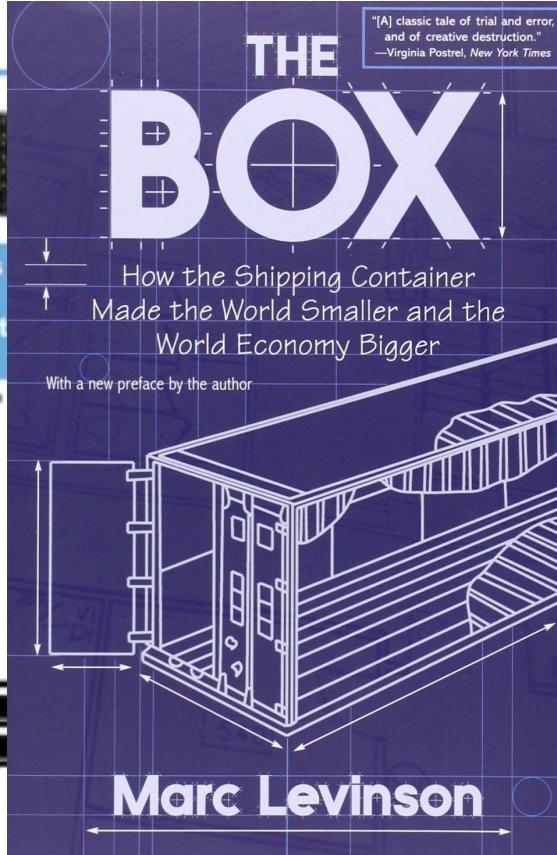
The shipping container

Multiple types of goods



A standard container that is loaded with virtually any goods, and stays sealed until it reaches final delivery.

Multiple methods of transportation



Do I worry about how goods interact? (i.e. place coffee beans next to spices)

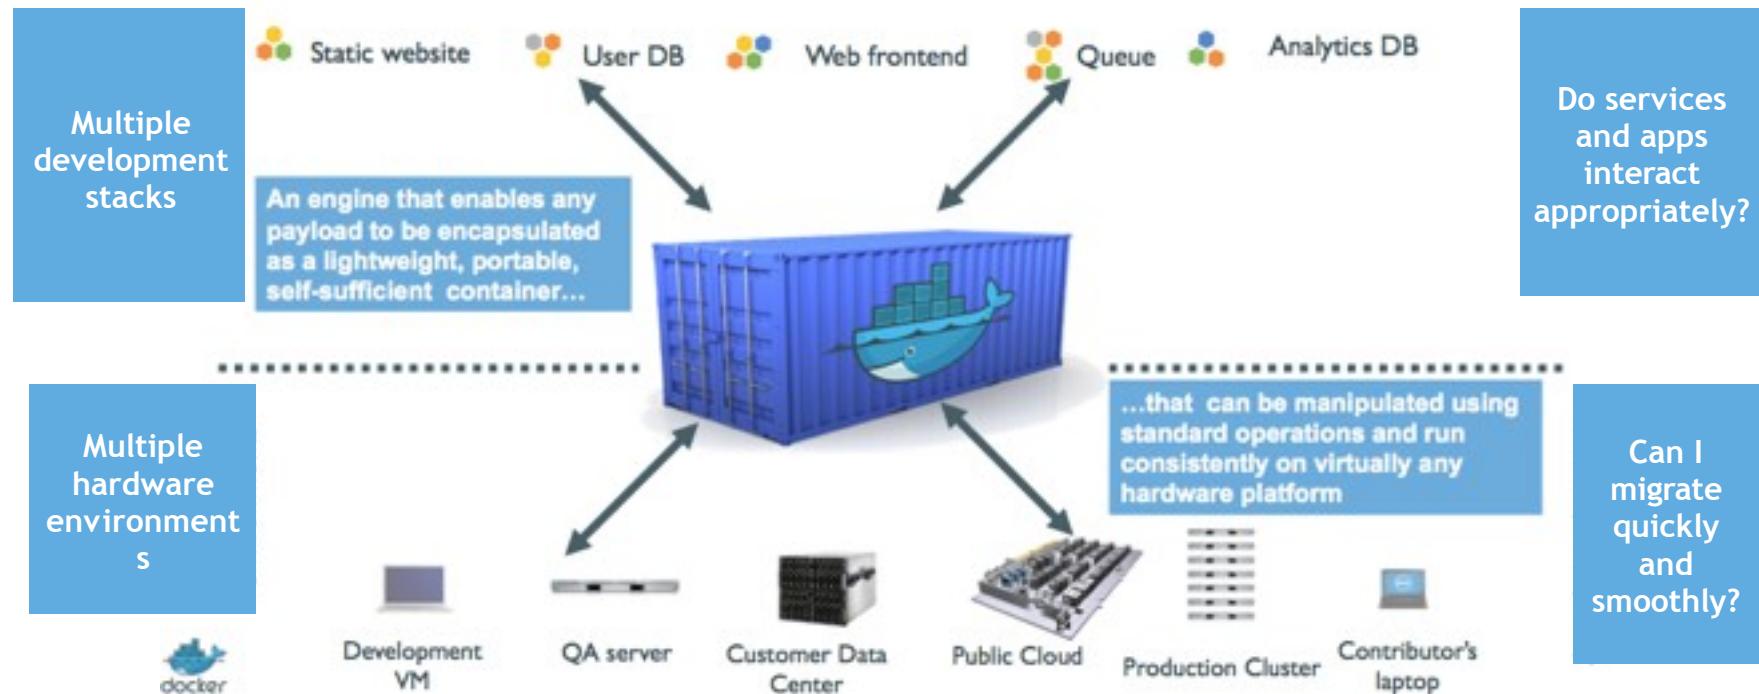
between, can be loaded and unloaded, stacked, transported quickly over long distances, transferred from one mode of transport to another



Can I transport quickly and smoothly? (i.e. unload from ship onto train)



Docker containers



Benefits of Docker

- Separation of concerns
 - Developers focus on building their apps
 - System admins focus on deployment
- Fast development cycle
- Application portability
 - Build in one environment, ship to another
- Scalability
 - Easily spin up new containers if needed
- Run more apps on one host machine



Module 2: Docker Concepts and Terms



Module objectives

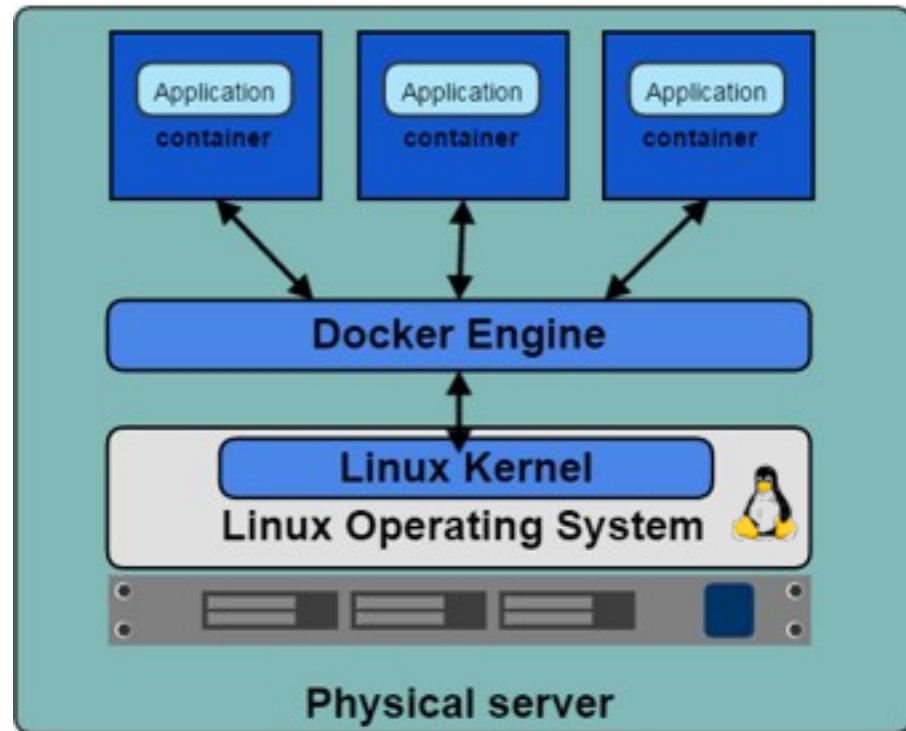
In this module we will:

- Explain all the main Docker concepts including
 - Docker Engine
 - Docker client
 - Containers
 - Images
 - Registry and Repositories
 - Docker Hub
 - Orchestration



Docker and the Linux Kernel

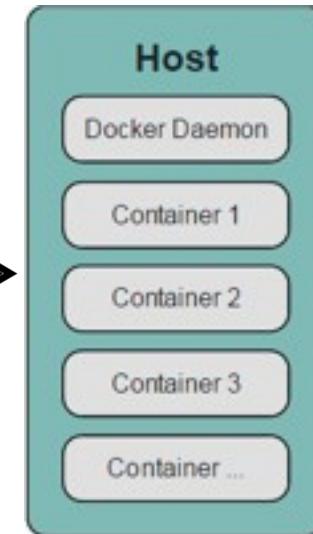
- **Docker Engine** is the program that enables containers to be distributed and run
- Docker Engine uses Linux Kernel namespaces and control groups
- Namespaces give us the isolated workspace



Docker Client and Daemon

- Client / Server architecture
- Client takes user inputs and sends them to the daemon
- Daemon runs and distributes containers
- Client and daemon can run on the same host or on different hosts
- CLI client and GUI (Kitematic)

Client



Checking Client and Daemon Version

- Run

```
docker version
```

```
1. vagrant@polytech: ~ (ssh)
root@polytech:~# docker version
Client:
  Version:          18.09.5
  API version:     1.39
  Go version:       go1.10.8
  Git commit:       e8ff056dbc
  Built:            Thu Apr 11 04:44:28 2019
  OS/Arch:          linux/amd64
  Experimental:    false

Server: Docker Engine - Community
Engine:
  Version:          18.09.5
  API version:     1.39 (minimum version 1.12)
  Go version:       go1.10.8
  Git commit:       e8ff056
  Built:            Thu Apr 11 04:10:53 2019
  OS/Arch:          linux/amd64
  Experimental:    false
```

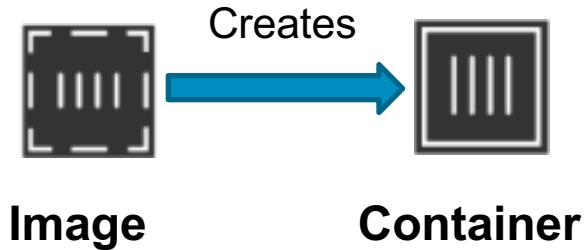
Client

Daemon

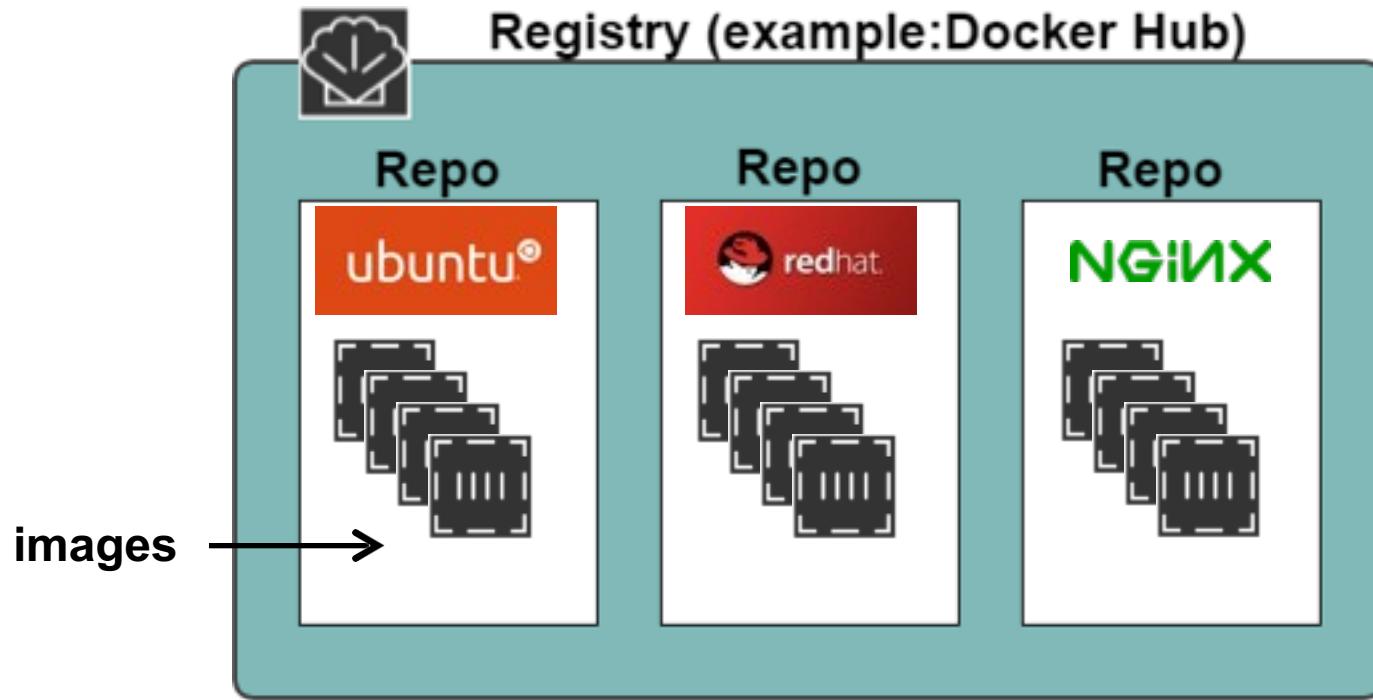


Docker Containers and Images

- **Images**
 - Read only template used to create containers (Binaries)
 - Built by you or other Docker users
 - Stored in Docker Hub, Docker Trusted Registry or your own Registry
- **Containers**
 - Isolated application platform
 - Contains everything needed to run your application
 - Based on one or more images

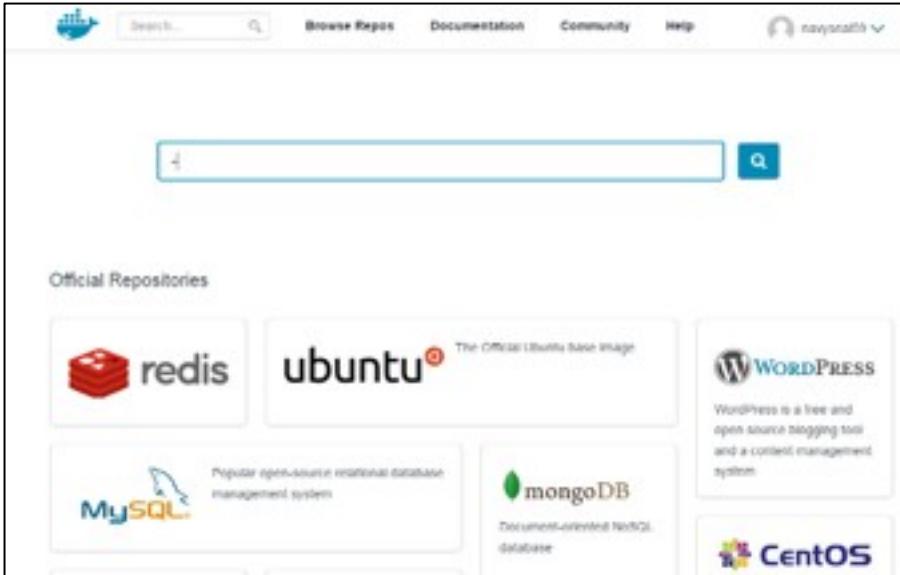


Registry and Repository



Docker Hub

Docker Hub is the public registry that contains a large number of images available for your use



Docker Orchestration

- Four tools for orchestrating distributed applications with Docker
- Docker Machine
 - Tool that provisions Docker hosts and installs the Docker Engine on them
- Docker Swarm
 - Tool that clusters many Engines and schedules containers
- Docker Compose
 - Tool to create and manage multi-container applications
- Kubernetes
 - Google Borg for community (swarm killer)



Module 3: Installing Docker



Module objectives

In this module we will:

- Outline the ways to install Docker on various Linux operating systems
- Install Docker on our Amazon AWS Ubuntu instance
- Explain how to run Docker without requiring “sudo”
- Install Boot2Docker on our PC or Mac



Installation options overview

- Docker can be installed in a variety of ways depending on your platform
- Packaged distributions available for most Linux platforms
 - Ubuntu, CentOS, Fedora, Arch Linux, RHEL, Debian, Gentoo, openSUSE
- Binary download
- Installation script from Docker



Module 4: Introduction to Images



Module objectives

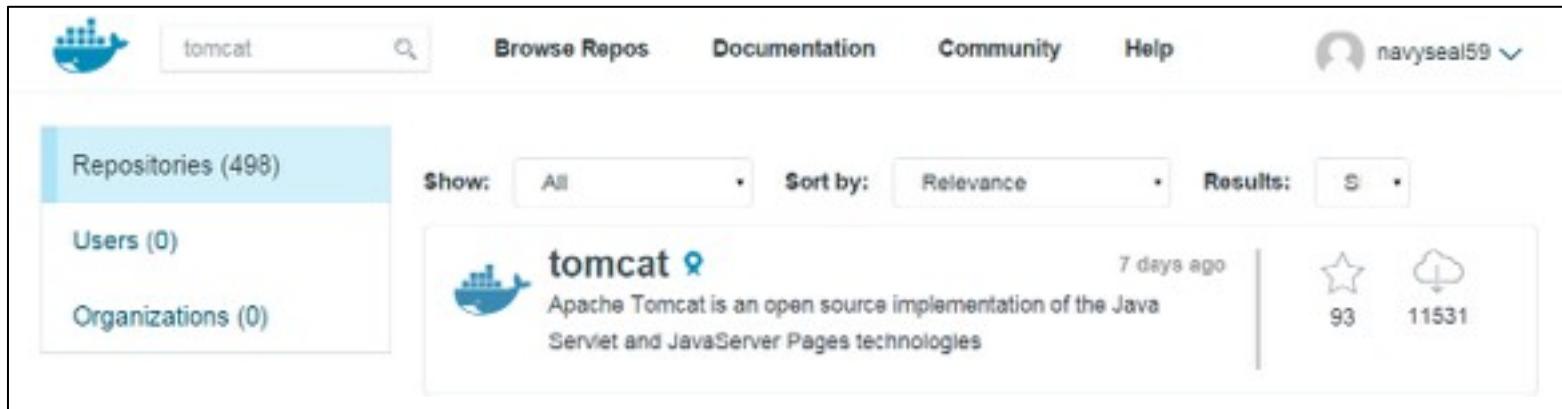
In this module we will:

- Learn to search for images on Docker Hub and with the Docker client
- Explain what official repositories are
- Create a Docker Hub account
- Explain the concept of image tags
- Learn to pull images from Docker Hub



Search for Images on Docker Hub

- Lots of Images available for use
- Images reside in various Repositories



The screenshot shows the Docker Hub search interface. The search bar at the top contains the text "tomcat". Below the search bar, there are navigation links: "Browse Repos", "Documentation", "Community", and "Help". On the right, there is a user profile for "navyseal59". The main search results are displayed in a table. The first result is for the "tomcat" repository, which has 498 repositories. The repository card for "tomcat" shows the Docker logo, the name "tomcat" with a blue star icon, and the description "Apache Tomcat is an open source implementation of the Java Servlet and JavaServer Pages technologies". It also shows the last update time as "7 days ago" and the statistics "93" for stars and "11531" for downloads.

Repositories (498)
Users (0)
Organizations (0)

tomcat

Apache Tomcat is an open source implementation of the Java Servlet and JavaServer Pages technologies

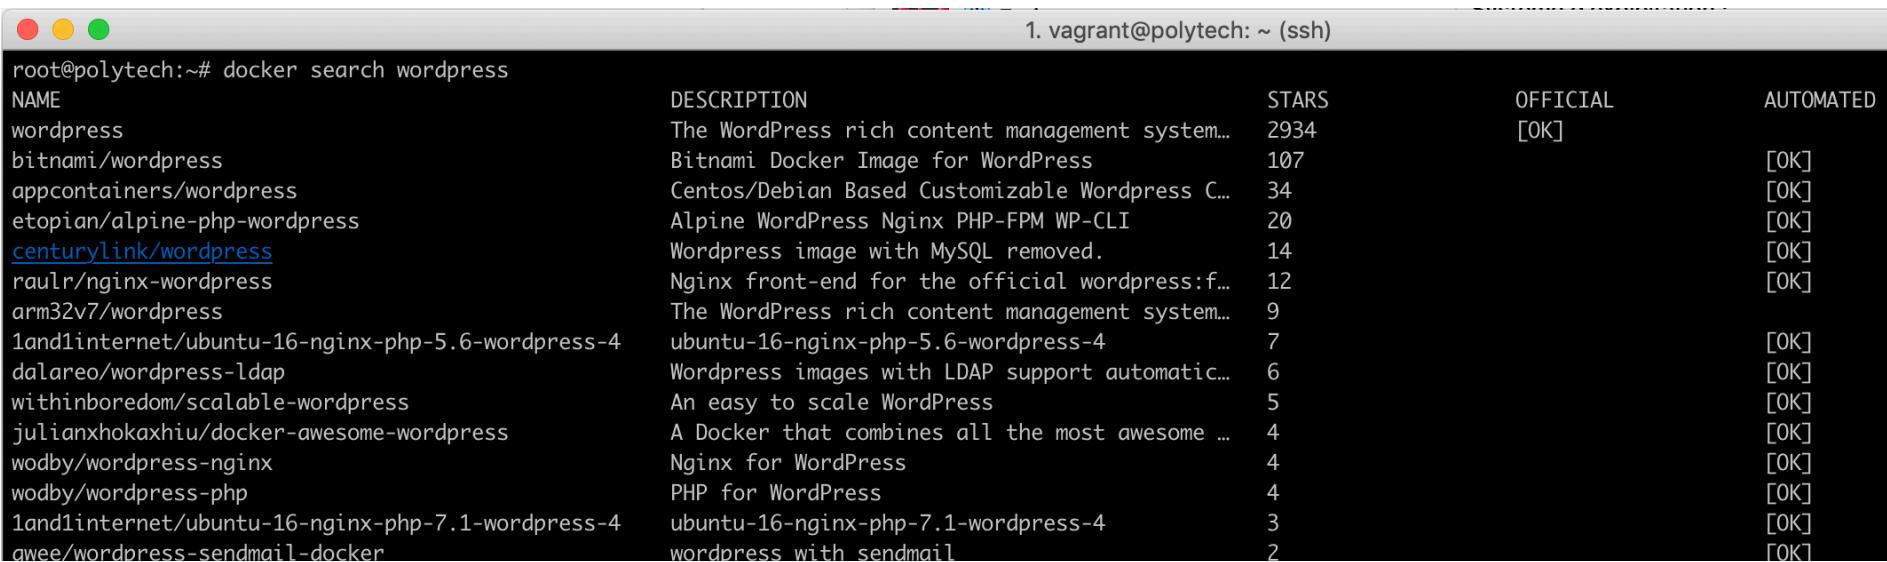
7 days ago

93 11531



Search for images using Docker client

- Run the docker search command
- Results displayed in table form



1. vagrant@polytech: ~ (ssh)

NAME	DESCRIPTION	STARS	OFFICIAL	AUTOMATED
wordpress	The WordPress rich content management system...	2934	[OK]	
bitnami/wordpress	Bitnami Docker Image for WordPress	107	[OK]	
appcontainers/wordpress	Centos/Debian Based Customizable Wordpress C...	34	[OK]	
etopian/alpine-php-wordpress	Alpine WordPress Nginx PHP-FPM WP-CLI	20	[OK]	
centurylink/wordpress	Wordpress image with MySQL removed.	14	[OK]	
raulr/nginx-wordpress	Nginx front-end for the official wordpress:f...	12	[OK]	
arm32v7/wordpress	The WordPress rich content management system...	9		
1and1internet/ubuntu-16-nginx-php-5.6-wordpress-4	ubuntu-16-nginx-php-5.6-wordpress-4	7	[OK]	
dalareo/wordpress-ldap	Wordpress images with LDAP support automatic...	6	[OK]	
withinboredom/scalable-wordpress	An easy to scale WordPress	5	[OK]	
julianxhokaxhiu/docker-awesome-wordpress	A Docker that combines all the most awesome ...	4	[OK]	
wodby/wordpress-nginx	Nginx for WordPress	4	[OK]	
wodby/wordpress-php	PHP for WordPress	4	[OK]	
1and1internet/ubuntu-16-nginx-php-7.1-wordpress-4	ubuntu-16-nginx-php-7.1-wordpress-4	3	[OK]	
aewe/wordpress-sendmail-docker	wordpress with sendmail	2	[OK]	



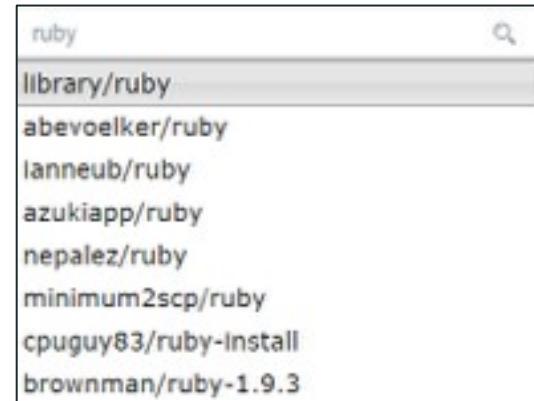
Official repositories

- Official repositories are a certified and curated set of Docker repositories that are promoted on Docker Hub
- Repositories come from vendors such as NGINX, Ubuntu, Red Hat, Redis, etc...
- Images are **supported by their maintainers**, optimised and up to date
- Official repository images are a mixture of
 - Base images for Linux operating systems (Ubuntu, CentOS etc...)
 - Images for popular development tools, programming languages, web and application servers, data stores



Identifying an official repository

- There are a few ways to tell if a repository is official
 - Marked on the OFFICIAL column in the terminal output
 - Prefixed with “library” on autocomplete search results in Docker Hub
 - Repository has the Docker logo on the Docker Hub search results



A screenshot of a Docker Hub repository page for the 'ruby' repository. The page includes search and filter controls at the top: 'Show: All', 'Sort by: Relevance', and 'Results: S'. The main content area shows the repository details: a blue whale icon, the repository name 'ruby' in bold, a question mark icon, and a description: 'Ruby is a dynamic, reflective, object-oriented, general-purpose, open-source programming language.' Below this, a timestamp '4 days ago' is shown. To the right, there are two icons: a star icon with the number '164' and a cloud icon with the number '148540'.



Display Local Images

- Run docker images
- When creating a container Docker will attempt to use a local image first
- If no local image is found, the Docker daemon will look in Docker Hub unless another registry is specified



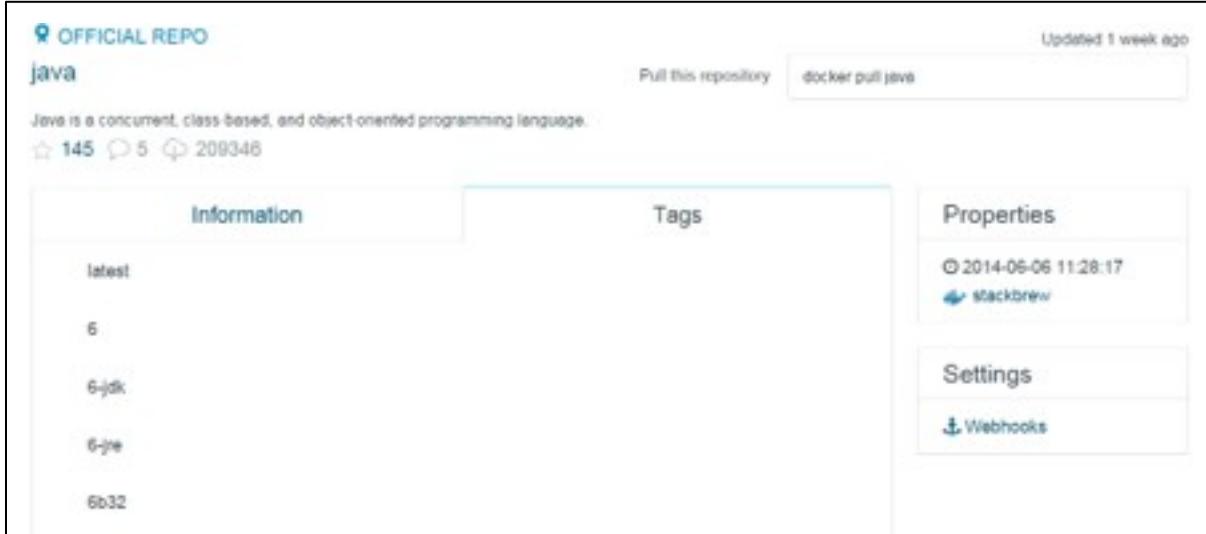
REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
nginx	latest	27a188018e18	11 days ago	109MB
wordpress	latest	837092bc87de	2 weeks ago	421MB
hello-world	latest	fce289e99eb9	3 months ago	1.84kB

root@polytech:~#



Image Tags

- Images are specified by **repository:tag**
- The same image may have multiple tags
- The default tag is `latest`
- Look up the repository on Docker Hub to see what tags are available



The screenshot shows the Docker Hub page for the 'java' repository. At the top, it says 'OFFICIAL REPO' and 'java'. Below that, a description states 'Java is a concurrent, class-based, and object-oriented programming language.' with statistics: 145 stars, 5 forks, and 209346 releases. A 'Pull this repository' button and a 'docker pull java' command are also present. The page is last updated '1 week ago'. The main content area is divided into three tabs: 'Information', 'Tags', and 'Properties'. The 'Tags' tab is selected, showing a list of available tags: 'latest', '6', '6-jdk', '6-jre', and '6b32'. The 'Properties' tab shows the creation date '2014-06-06 11:28:17' and the maintainer 'stackbrew'. The 'Settings' tab is partially visible.



Pulling images

- To download an image from Docker Hub or any registry, use `docker pull` command
- When running a container with the `docker run` command, images are automatically pulled if no local copy is found

Pull the latest image from the Ubuntu repository in Docker Hub

```
docker pull ubuntu
```

Pull the image with tag 12.04 from Ubuntu repository in Docker Hub

```
docker pull ubuntu:12.04
```



Module 5:

Running and Managing Containers



Module objectives

In this module we will

- Learn how to launch containers with the `docker run` command
- Explore different methods of accessing a container
- Explain how container processes work
- Learn how to stop and start containers
- Learn how to check container logs
- Learn how to find your containers with the `docker ps` command
- Learn how to display information about a container



Container lifecycle

- Basic lifecycle of a Docker container
 - Create container from image
 - Run container with a specified process
 - Process finishes and container stops
 - Destroy container
- More advanced lifecycle
 - Create container from image
 - Run container with a specified process
 - Interact and perform other actions inside the container
 - Stop the container
 - Restart the container



Creating and running a Container

- Use docker run command
- The docker run command actually does two things
 - Creates the container using the image we specify
 - Runs the container
- Syntax
docker run [options] [image] [command] [args]
- Image is specified with repository:tag

Examples

```
docker run ubuntu:14.04 echo "Hello World"
```

```
docker run ubuntu ps ax
```



Find your Containers

- Use `docker ps` to list running containers
- The `-a` flag to list all containers (includes containers that are stopped)



1. vagrant@polytech: ~ (ssh)

```
root@polytech:~# docker ps
CONTAINER ID        IMAGE               COMMAND
root@polytech:~# docker ps -a
CONTAINER ID        IMAGE               COMMAND
d5a122654320        debian              "bash"
c966c7ad66ec        wordpress          "docker-entrypoint.s..."
f710b9acecef        hello-world        "/hello"
c5c0e162308a        hello-world        "/hello"
e5b550c10543        hello-world        "/hello"
c6b49b6c611f        hello-world        "/hello"
39cce39c990b        hello-world        "/hello"
d5e1006c9fea        nginx              "nginx -g 'daemon of...'"
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
d5a122654320	debian	"bash"	5 minutes ago	Exited (0) 5 minutes ago		
c966c7ad66ec	wordpress	"docker-entrypoint.s..."	13 minutes ago	Exited (0) 11 minutes ago		
f710b9acecef	hello-world	"/hello"	28 minutes ago	Exited (0) 28 minutes ago		
c5c0e162308a	hello-world	"/hello"	28 minutes ago	Exited (0) 28 minutes ago		
e5b550c10543	hello-world	"/hello"	28 minutes ago	Exited (0) 28 minutes ago		
c6b49b6c611f	hello-world	"/hello"	28 minutes ago	Exited (0) 28 minutes ago		
39cce39c990b	hello-world	"/hello"	30 minutes ago	Exited (0) 30 minutes ago		
d5e1006c9fea	nginx	"nginx -g 'daemon of..."'	30 minutes ago	Exited (0) 25 minutes ago		

```
root@polytech:~#
```



Container with Terminal

- Use `-i` and `-t` flags with `docker run`
- The `-i` flag tells docker to connect to STDIN on the container
- The `-t` flag specifies to get a pseudo-terminal
- **Note:** You need to run a terminal process as your command (e.g. `bash`)

Example

```
docker run -i -t ubuntu:latest bash
```



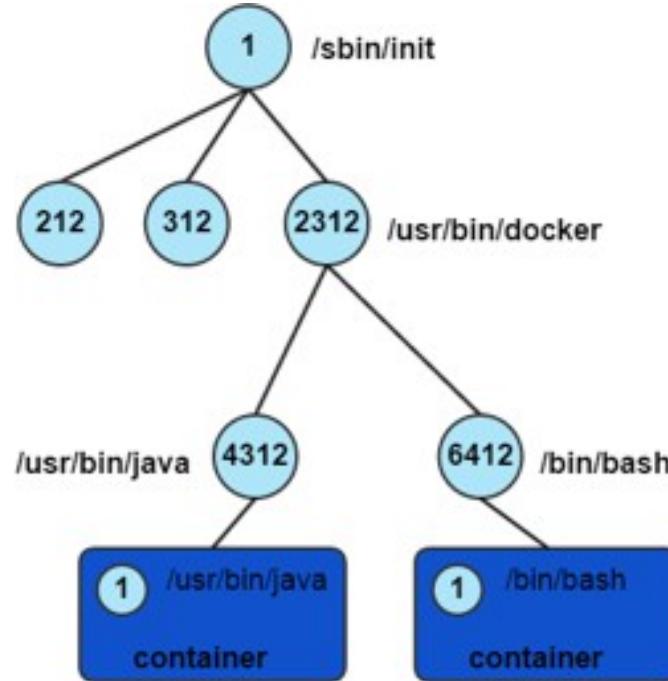
Exit the Terminal

- Type `exit` to quit the terminal and return to your host terminal
- Exiting the terminal will shutdown the container
- To exit the terminal without a shutdown, hit `CTRL + P + Q` together



Container Processes

- A container only runs as long as the process from your specified `docker run` command is running
- Your command's process is always PID 1 inside the container



Container ID

- Containers can be specified using their ID or name
- Long ID and short ID
- Short ID and name can be obtained using `docker ps` command to list containers
- Long ID obtained by inspecting a container



docker ps command

- To view only the container ID's (displays short ID)

```
docker ps -q
```

- To view the last container that was started

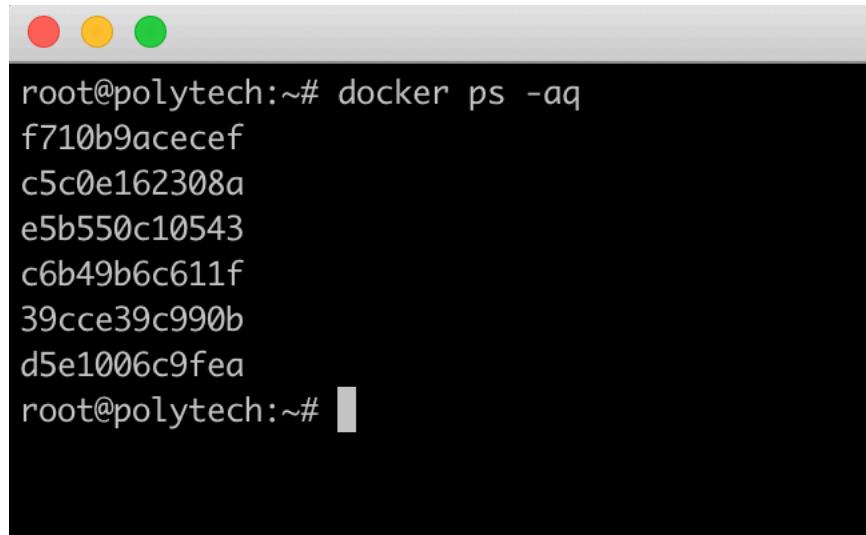
```
docker ps -l
```

```
1. vagrant@polytech: ~ (ssh)
root@polytech:~# docker ps -l
CONTAINER ID        IMAGE               COMMAND             CREATED            STATUS              PORTS               NAMES
f710b9acecef        hello-world        "/hello"           49 seconds ago   Exited (0) 48 seconds ago
root@polytech:~#
```



docker ps command

- Combining flags to list all containers with only their short ID
docker ps -aq
- Combining flags to list the short ID of the last container started
docker ps -lq

A screenshot of a terminal window on a Mac OS X desktop. The window has the standard red, yellow, and green title bar buttons. The terminal itself is black with white text. The text shows the command 'root@polytech:~# docker ps -aq' followed by a list of container IDs: f710b9acecef, c5c0e162308a, e5b550c10543, c6b49b6c611f, 39cce39c990b, and d5e1006c9fea. The prompt 'root@polytech:~#' is visible at the bottom.

```
root@polytech:~# docker ps -aq
f710b9acecef
c5c0e162308a
e5b550c10543
c6b49b6c611f
39cce39c990b
d5e1006c9fea
root@polytech:~#
```



docker ps filtering

- Use the `--filter` flag to specify filtering conditions
- Currently you can filter based on the container's exit code and status
- Status can be one of
 - Restarting
 - Running
 - Exited
 - Paused
- To specify multiple conditions, pass multiple `--filter` flags



docker ps filtering examples

- List containers with an exit code of 1 (exited with error)

```
1. vagrant@polytech: ~ (ssh)
root@polytech:~# docker ps -a --filter "exited=0"
CONTAINER ID        IMAGE               COMMAND             CREATED            STATUS              PORTS               NAMES
39cce39c990b        hello-world        "/hello"           About a minute ago   Exited (0) About a minute ago
root@polytech:~#
```



Running in Detached Mode

- Also known as running in the background or as a daemon
- Use `-d` flag
- To observe output use `docker logs [container id]`

Create a centos container and run the ping command to ping the container itself 50 times

```
docker run -d centos:7 ping 127.0.0.1 -c 50
```



A More Practical Container

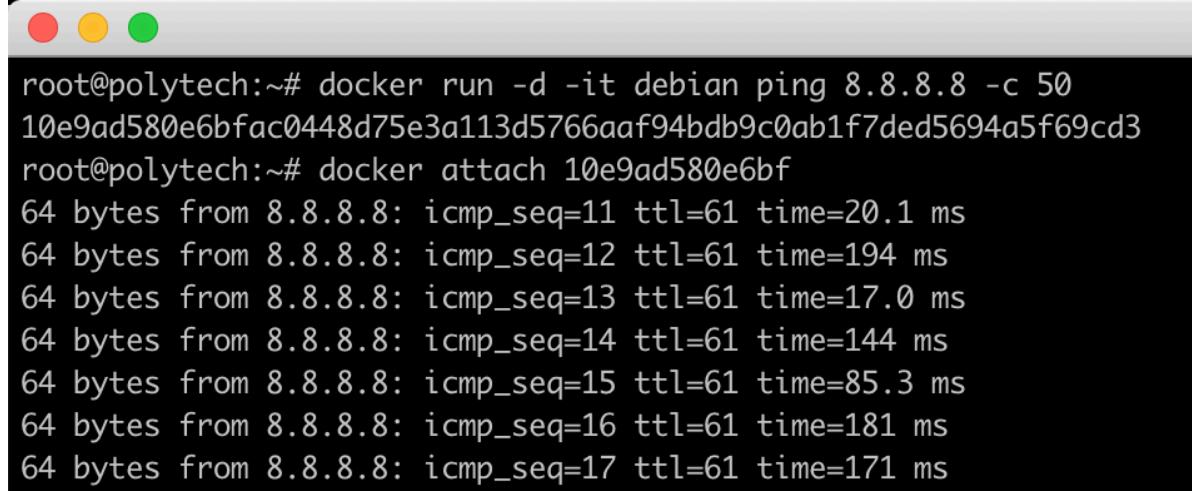
- Run a web application inside a container
- The `-P` flag to map container ports to host ports

Create a container using the tomcat image, run in detached mode and map the tomcat ports to the host port
`docker run -d -P tomcat:7`



Attaching to a container

- Attaching a client to a container will bring a container which is running in the background into the foreground
- The containers PID 1 process output will be displayed on your terminal
- Use `docker attach` command and specify the container ID or name
- **Warning:** Attaching to containers is error prone because if you hit `CTRL + C` by accident, you will stop the process and therefore stop the container

A screenshot of a terminal window with a light gray header bar containing three colored window control buttons (red, yellow, green). The main area of the terminal is black with white text. It shows the following command and its output:

```
root@polytech:~# docker run -d -it debian ping 8.8.8.8 -c 50
10e9ad580e6bfac0448d75e3a113d5766aaf94bdb9c0ab1f7ded5694a5f69cd3
root@polytech:~# docker attach 10e9ad580e6bf
64 bytes from 8.8.8.8: icmp_seq=11 ttl=61 time=20.1 ms
64 bytes from 8.8.8.8: icmp_seq=12 ttl=61 time=194 ms
64 bytes from 8.8.8.8: icmp_seq=13 ttl=61 time=17.0 ms
64 bytes from 8.8.8.8: icmp_seq=14 ttl=61 time=144 ms
64 bytes from 8.8.8.8: icmp_seq=15 ttl=61 time=85.3 ms
64 bytes from 8.8.8.8: icmp_seq=16 ttl=61 time=181 ms
64 bytes from 8.8.8.8: icmp_seq=17 ttl=61 time=171 ms
```



Detaching from a container

- Hit CTRL + P + Q together on your terminal
- Only works if the following two conditions are met
 - The container standard input is connected
 - The container has been started with a terminal
 - For example: `docker run -i -t ubuntu`
- Hitting CTRL + C will terminate the process, thus shutting down the container



Docker exec command

- `docker exec` command allows us to execute additional processes inside a container
- Typically used to gain command line access
- `docker exec -i -t [container ID] bash`
- Exiting from the terminal will not terminate the container



Inspecting container logs

- Container PID 1 process output can be viewed with `docker logs` command
- Will show whatever PID 1 writes to stdout
- Displays the entire log output from the time the container was created

View the output of the containers PID 1 process

```
docker logs <container name>
```



Following container logs

- Use `docker logs` command and specify the `-f` option
- Similar to Linux `tail -f` command
- `CTRL + C` to exit
- Will still follow the log output from the very beginning



Tailing container logs

- We can specify to only show the last “x” number of lines from the logs
- Use --tail option and specify the number

Show the last 5 lines from the container log

```
docker logs --tail 5 <container ID>
```

Show the last 5 lines and follow the log

```
docker logs --tail 5 -f <container ID>
```



Stopping a container

- Two commands we can use
 - `docker stop`
 - `docker kill`
- `docker stop` sends a **SIGTERM** to the main container process
 - Process then receives a **SIGKILL** after a grace period
 - Grace period can be specified with `-t` flag (default is 10 seconds)
- `docker kill` sends a **SIGKILL** immediately to the main container process



Restarting a container

- Use `docker start` to restart a container that has been stopped
- Container will start using the same options and command specified previously
- Can attach to the container with `-a` flag

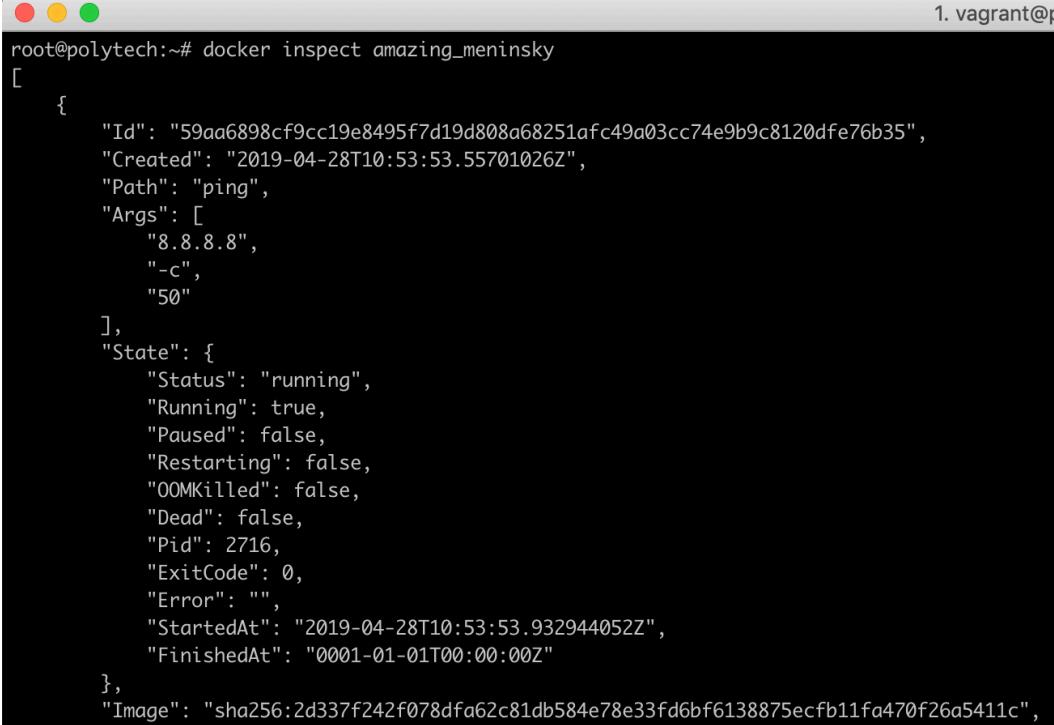
Start a stopped container and attach to the process that it is running

```
docker start -a <container ID>
```



Inspecting a container

- `docker inspect` command displays all the details about a container
- Outputs details in JSON array



```
root@polytech:~# docker inspect amazing_meninsky
[{"Id": "59aa6898cf9cc19e8495f7d19d808a68251afc49a03cc74e9b9c8120dfe76b35", "Created": "2019-04-28T10:53:53.55701026Z", "Path": "ping", "Args": ["8.8.8.8", "-c", "50"], "State": {"Status": "running", "Running": true, "Paused": false, "Restarting": false, "OOMKilled": false, "Dead": false, "Pid": 2716, "ExitCode": 0, "Error": "", "StartedAt": "2019-04-28T10:53:53.932944052Z", "FinishedAt": "0001-01-01T00:00:00Z"}, "Image": "sha256:2d337f242f078dfa62c81db584e78e33fd6bf6138875ecfb11fa470f26a5411c"}]
```



Finding a specific property

- You can pipe the output of `docker inspect` to `grep` and use it to search a particular container property.
- Example

```
docker inspect <container name> | grep IPAddress
```

```
johnnytu@docker-ubuntu:~$ docker inspect jolly_davinci | grep IPAddress
    "IPAddress": "172.17.0.24",
```



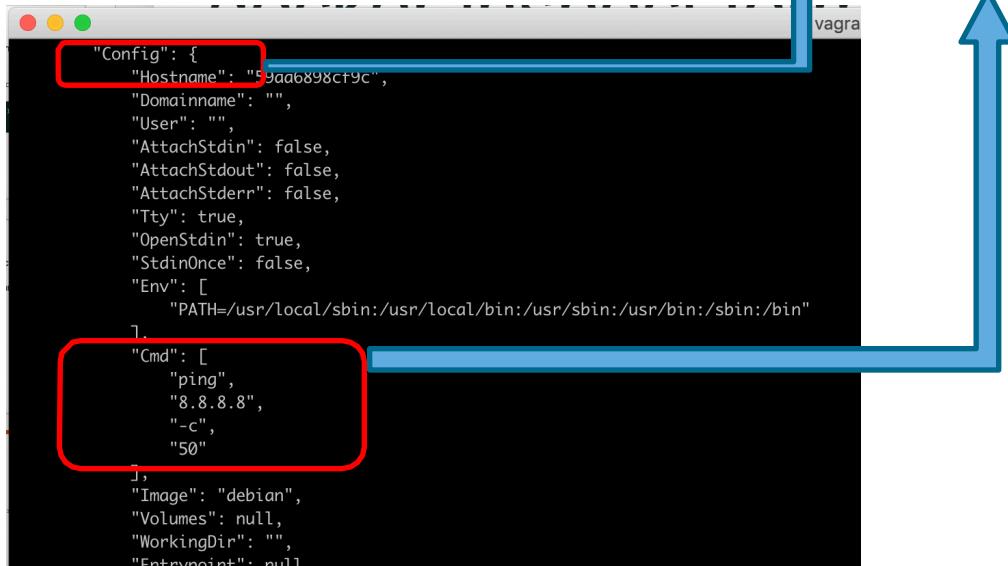
Formatting docker inspect output

- Piping the output to `grep` is simple, but not the most effective method of formatting
- Use the `--format` option of the `docker inspect` command
- Format option uses Go's text/template package
<http://golang.org/pkg/text/template/>



docker inspect formatting syntax

- docker inspect --format='{{.<field1>.<field2>}}' \<container id>
- Field names are case sensitive
- Example – to get the value of the Cmd field
docker inspect --format='{{.Config.Cmd}}' <container id>



```
vagra
{
  "Config": {
    "Hostname": "9aab898c19c",
    "Domainname": "",
    "User": "",
    "AttachStdin": false,
    "AttachStdout": false,
    "AttachStderr": false,
    "Tty": true,
    "OpenStdin": true,
    "StdinOnce": false,
    "Env": [
      "PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin"
    ],
    "Cmd": [
      "ping",
      "8.8.8.8",
      "-c",
      "50"
    ],
    "Image": "debian",
    "Volumes": null,
    "WorkingDir": "",
    "Entrypoint": null
  }
}
```



Inspecting a whole JSON object

- When you want to output all the fields of a JSON object you need to use the Go template's JSON function

```
root@polytech:~# docker inspect --format='{{json .Config}}' amazing_meninsky | jq
{
  "Hostname": "59aa6898cf9c",
  "Domainname": "",
  "User": "",
  "AttachStdin": false,
  "AttachStdout": false,
  "AttachStderr": false,
  "Tty": true,
  "OpenStdin": true,
  "StdinOnce": false,
  "Env": [
    "PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin"
  ],
  "Cmd": [
    "ping",
    "8.8.8.8",
    "-c",
    "50"
  ],
  "Image": "debian",
  "Volumes": null,
  "WorkingDir": "",
  "Entrypoint": null,
  "OnBuild": null,
  "Labels": {}
}
```

Incorrect approach

```
docker inspect --format='{{.Config}}' <container name>
```

Correct approach

```
docker inspect --format='{{json .Config}}' <container name>
```



Deleting containers

- Can only delete containers that have been stopped
- Use `docker rm` command
- Specify the container ID or name



Delete all containers

- Use `docker ps -aq` to list the id's of all containers
- Feed the output into `docker rm` command

Delete all containers that are stopped

```
docker rm $(docker ps -aq)
```



Module Summary

- Containers can be run in the foreground and background
- A container only runs as long as the process we specified during creation is running
- The container log is the output of its PID 1 process
- Key commands we learnt
 - docker run
 - docker ps
 - docker logs
 - docker inspect



Module 6: Building Images



Module objectives

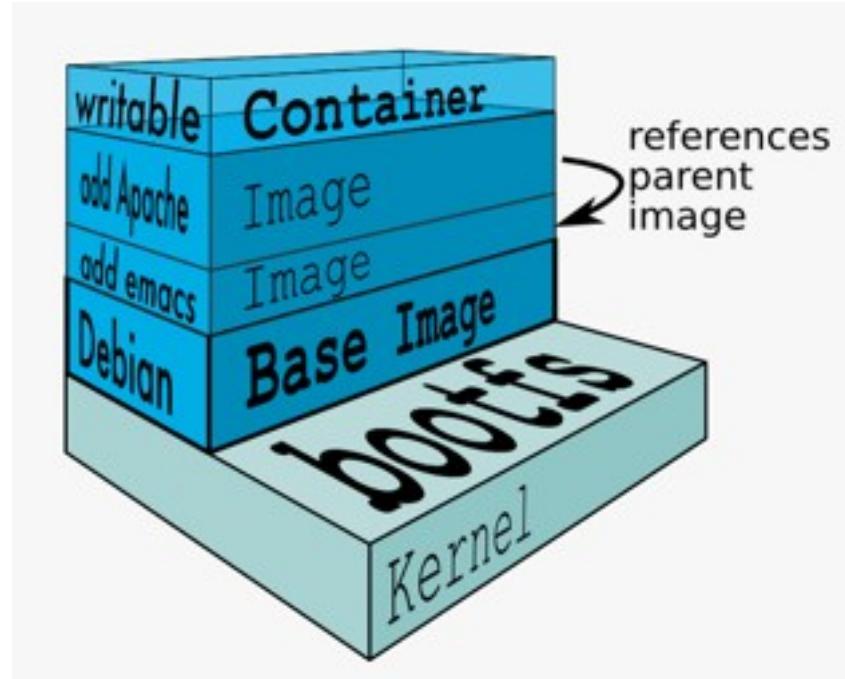
In this module we will

- Explain how image layers work
- Build an image by committing changes in a container
- Learn how to build images with a Dockerfile
- Work through examples of key Dockerfile instructions
 - RUN
 - CMD
 - ENTRYPOINT
 - COPY
- Talk about the best practices when writing a Dockerfile



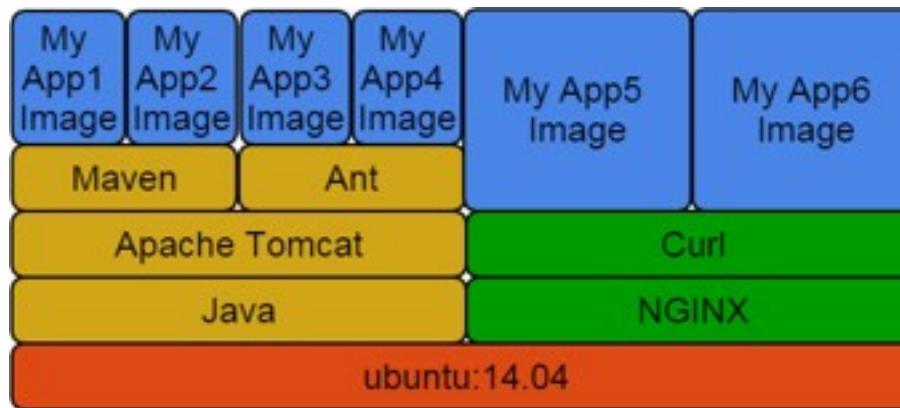
Understanding image layers

- An image is a collection of files and some meta data
- Images are comprised of multiple layers
- A layer is also just another image
- Each image contains software you want to run
- Every image contains a base layer
- Docker uses a copy on write system
- Layers are read only
- COW/Union Filesystems (AUFS/BTRFS)



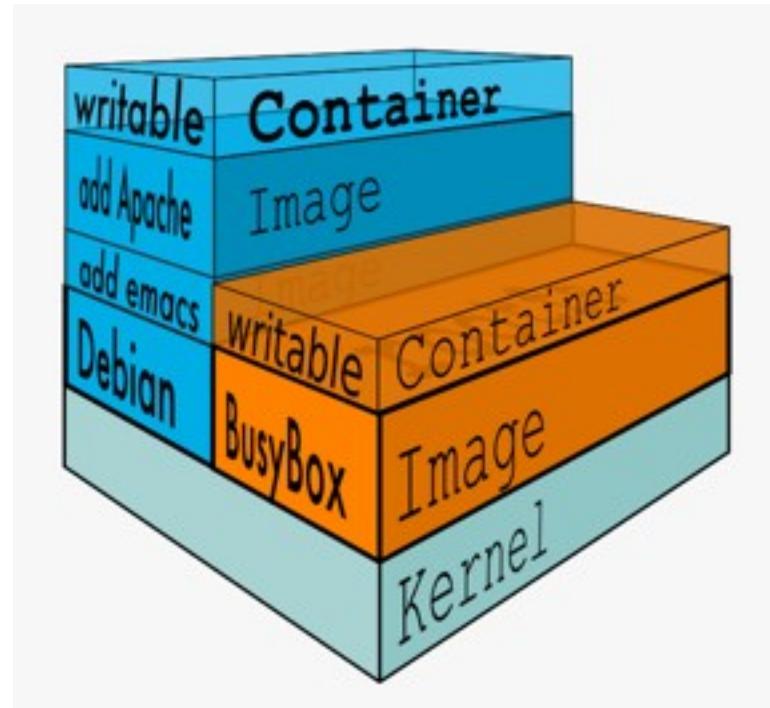
Sharing layers

- Images can share layers in order to speed up transfer times and optimize disk and memory usage
- Parent images that already exists on the host do not have to be downloaded



The container writable layer

- Docker creates a top writable layer for containers
- Parent images are read only
- All changes are made at the writable layer
- When changing a file from a read only layer, the copy on write system will copy the file into the writable layer



Methods of building images

- Three ways
 - Commit changes from a container as a new image
 - Build from a Dockerfile
 - Import a tarball into Docker as a standalone base layer



Committing changes in a container

- Allows us to build images interactively
- Get terminal access inside a container and install the necessary programs and your application
- Then save the container as a new image using the docker commit command



Comparing container changes

- Use the `docker diff` command to compare a container with its parent image
 - Recall that images are read only and changes occur in a new layer
 - The parent image (the original) is being compared with the new layer
- Copy on write system ensures that starting a container from a large image does not result in a large copy operation
- Lists the files and directories that have changed



```
root@polytech:~# docker diff youthful_kilby | head -50
C /usr
C /usr/lib
A /usr/lib/ssl
A /usr/lib/ssl/certs
A /usr/lib/ssl/misc
A /usr/lib/ssl/misc/c_name
A /usr/lib/ssl/misc/tsget
A /usr/lib/ssl/misc/CA.pl
A /usr/lib/ssl/misc/CA.sh
A /usr/lib/ssl/misc/c_hash
A /usr/lib/ssl/misc/c_info
```



Docker Commit

- docker commit command saves changes in a container as a new image
- **Syntax**
docker commit [options] [container ID] [repository:tag]
- Repository name should be based on username/application
- Can reference the container with container name instead of ID

Save the container with ID of 984d25f537c5 as a new image in the repository johnnytu/myapplication. Tag the image as 1.0

```
docker commit 984d25f537c5 johnnytu/myapplication:1.0
```



Image namespaces

- Image repositories belong in one of three namespaces
 - Root
 - ubuntu:14.04
 - centos:7
 - nginx
 - User OR organization
 - johnnytu/myapp
 - mycompany/myapp
 - Self hosted
 - registry.mycompany.com:5000/my-image



Uses for namespaces

- Root namespace is for official repositories
- User and organization namespaces are for images you create and plan to distribute on Docker Hub
- Self-hosted namespace is for images you create and plan to distribute in your own registry server



Intro to Dockerfile

*A **Dockerfile** is a configuration file that contains instructions for building a Docker image*

- Provides a more effective way to build images compared to using docker commit
- Easily fits into your development workflow and your continuous integration and deployment process



Process for building images from Dockerfile

1. Create a Dockerfile in a new folder or in your existing application folder
2. Write the instructions for building the image
 - What programs to install
 - What base image to use
 - What command to run
3. Run `docker build` command to build an image from the Dockerfile



Dockerfile Instructions

- Instructions specify what to do when building the image
- **FROM** instruction specifies what the base image should be
- **RUN** instruction specifies a command to execute
- Comments start with “#”

```
#Example of a comment
FROM ubuntu:14.04
RUN apt-get install vim
RUN apt-get install curl
```



FROM instruction

- Must be the first instruction specified in the Dockerfile (not including comments)
- Can be specified multiple times to build multiple images
 - Each `FROM` marks the beginning of a new image
- Can use any image including, images from official repositories, user images and images in self hosted registries.

Examples

```
FROM ubuntu
FROM ubuntu:14.04
FROM johnnytu/myapplication:1.0
FROM company.registry:5000/myapplication:1.0
```



More about RUN

- RUN will do the following:
 - Execute a command.
 - Record changes made to the filesystem.
 - Works great to install libraries, packages, and various files.
- RUN will NOT do the following:
 - Record state of *processes*.
 - Automatically start daemons.



Docker Build

- **Syntax**

```
docker build [options] [path]
```

- **Common option to tag the build**

```
docker build -t [repository:tag] [path]
```

Build an image using the current folder as the context path. Put the image in the johnnytu/myimage repository and tag it as 1.0

```
docker build -t johnnytu/myimage:1.0 .
```

As above but use the myproject folder as the context path

```
docker build -t johnnytu/myimage:1.0 myproject
```



Build output

```
johnnytu@docker-ubuntu:~/myimage$ docker build -t myimage .
Sending build context to Docker daemon 2.048 kB
Sending build context to Docker daemon
Step 0 : FROM ubuntu:14.04
--> 07f8e8c5e660
Step 1 : RUN apt-get update
--> Running in fd009309d260
...
...
Reading package lists...
--> 161253fe4244
Removing intermediate container fd009309d260
Step 2 : RUN apt-get install -y wget
--> Running in 69ba8a082150
Reading package lists...
...
...
--> c7cc72b567e4
Removing intermediate container 69ba8a082150
Successfully built c7cc72b567e4
```



The build context

```
johnnytu@docker-ubuntu:~/myimage$ docker build -t myimage .
Sending build context to Docker daemon 2.048 kB
Sending build context to Docker daemon
```

- The build context is the directory that the Docker client sends to the Docker daemon during the `docker build` command
- Directory is sent as an archive
- Docker daemon will build using the files available in the context
- Specifying “.” for the build context means to use the current directory



Examining the build process

- Each RUN instruction will execute the command on the top writable layer of a new container

```
Step 1 : RUN apt-get update
--> Running in fd009309d260
```

- At the end of the execution of the command, the container is committed as a new image and then deleted

```
--> 161253fe4244
Removing intermediate container fd009309d260
```



Examining the build process (cont'd)

- The next RUN instruction will be executed in a new container using the newly created image

```
Step 2 : RUN apt-get install -y wget
--> Running in 69ba8a082150
```

- The container is committed as a new image and then removed. Since this RUN instruction is the final instruction, the image committed is what will be used when we specify the repository name and tag used in the docker build command

```
--> c7cc72b567e4
Removing intermediate container 69ba8a082150
Successfully built c7cc72b567e4
```



The build cache

- Docker saves a snapshot of the image after each build step
- Before executing a step, Docker checks to see if it has already run that build sequence previously
- If yes, Docker will use the result of that instead of executing the instruction again
- Docker uses exact strings in your Dockerfile to compare with the cache
 - Simply changing the order of instructions will invalidate the cache
- To disable the cache manually use the `--no-cache` flag
`docker build --no-cache myimage .`



Run Instruction aggregation (Best Practice)

- Can aggregate multiple RUN instructions by using “`&&`”
- Commands will all be run in the same container and committed as a new image at the end
- Reduces the number of image layers that are produced

```
RUN apt-get update && apt-get install -y \
    curl \
    vim \
    openjdk-7-jdk
```



Viewing Image layers and history

- `docker history` command shows us the layers that make up an image
- See when each layer was created, its size and the command that was run

IMAGE	CREATED	CREATED BY	SIZE
10f1e1747aa1	12 seconds ago	/bin/sh -c apt-get install -y wget	6.119 MB
9b6aeeff1e9cc	23 seconds ago	/bin/sh -c apt-get install -y vim	43.12 MB
334d0289feff	10 minutes ago	/bin/sh -c apt-get update	20.86 MB
07f8e8c5e660	2 weeks ago	/bin/sh -c #(nop) CMD ["/bin/bash"]	0 B
37bea4ee0c81	2 weeks ago	/bin/sh -c sed -i 's/^#\s*/(deb.*universe)\\$/	1.895 kB
a82efea989f9	2 weeks ago	/bin/sh -c echo '#!/bin/sh' > /usr/sbin/polic	194.5 kB
e9e06b06e14c	2 weeks ago	/bin/sh -c #(nop) ADD file:f4d7b4b3402b5c53f2	188.1 MB



CMD Instruction

- CMD defines a default command to execute when a container is created
- Shell format and EXEC format
- Can only be specified once in a Dockerfile
 - If specified multiple times, the last CMD instruction is executed
- **Can be overridden at run time**
- (Array format doesn't append /bin/sh -c)

Shell format

```
CMD ping 127.0.0.1 -c 30
```

Exec format

```
CMD ["ping", "127.0.0.1", "-c", "30"]
```



ENTRYPOINT Instruction

- Defines the command that will run when a container is executed
- Run time arguments and CMD instruction are passed as parameters to the ENTRYPOINT instruction
- Shell and EXEC form
- Container essentially runs as an executable
- By default ENTRYPOINT is not overridden (can be using —entrypoint flag)

```
ENTRYPOINT ["/usr/sbin/nginx"]
CMD ["-h"]
```



Using CMD with ENTRYPOINT

- If ENTRYPOINT is used, the CMD instruction can be used to specify default parameters
- Parameters specified during docker run will override CMD
- If no parameters are specified during docker run, the CMD arguments will be used for the ENTRYPOINT command



Shell vs exec format

- The RUN, CMD and ENTRYPOINT instructions can be specified in either shell or exec form

In shell form, the command will run inside a shell with /bin/sh -c

```
RUN apt-get update
```

Exec format allows execution of command in images that don't have /bin/sh

```
RUN ["apt-get", "update"]
```



Shell vs exec format

- Shell form is easier to write and you can perform shell parsing of variables
- For example

```
CMD sudo -u $(USER) java ....
```
- Exec form does not require image to have a shell
- For the ENTRYPOINT instruction, using shell form will prevent the ability to specify arguments at run time
 - The CMD arguments will not be used as parameters for ENTRYPOINT



Overriding ENTRYPOINT

- Specifying parameters during `docker run` will result in your parameters being used as arguments for the `ENTRYPOINT` command
- To override the command specified by `ENTRYPOINT`, use the `--entrypoint` flag.
- Useful for troubleshooting your images

Run a container using the image “myimage” and specify to run a bash terminal instead of the program specified in the image `ENTRYPOINT` instruction

```
docker run -it --entrypoint bash myimage
```



Copying source files

- When building “real” images you would want to do more than just install some programs
- Examples
 - Compile your source code and run your application
 - Copy configuration files
 - Copy other content
- How do we get our content on our host into the container?
- Use the `COPY` instruction



COPY instruction

- The `COPY` instruction copies new files or directories from a specified **source** and adds them to the container filesystem at a specified **destination**
- Syntax
`COPY <src> <dest>`
- The `<src>` path must be inside the build context
- If the `<src>` path is a directory, all files in the directory are copied. The directory itself is not copied
- You can specify multiple `<src>` directories



COPY examples

Copy the server.conf file in the build context into the root folder of the container

```
COPY server.conf /
```

Copy the files inside the data/server folder of the build context into the /data/server folder of the container

```
COPY data/server /data/server
```



Dockerize an application

- The Dockerfile is essential if we want to adapt our existing application to run on containers
- Take a simple Java program as an example. To build and run it, we need the following on our host
 - The Java Development Kit (JDK)
 - The Java Virtual Machine (JVM)
 - Third party libraries depending on the application itself
- You compile the code, run the application and everything looks good



Dockerize an application

- Then you distribute the application and run it on a different environment and it fails
- Reasons why the Java application fails?
 - Missing libraries in the environment
 - Missing the JDK or JVM
 - Wrong version of libraries
 - Wrong version of JDK or JVM
- So why not run your application in a Docker container?
- Install all the necessary libraries in the container
- Build and run the application inside the container and distribute the image for the container
- Will run on any environment with the Docker Engine installed



Specify a working directory

- Previously all our instructions have been executed at the root folder in our container
- `WORKDIR` instruction allows us to set the working directory for any subsequent `RUN`, `CMD`, `ENTRYPOINT` and `COPY` instructions to be executed in
- **Syntax**
`WORKDIR /path/to/folder`
- Path can be absolute or relative to the current working directory
- Instruction can be used multiple times



MAINTAINER Instruction

- Specifies who wrote the Dockerfile
- Optional but best practice to include
- Usually placed straight after the FROM instruction

Example

```
MAINTAINER Docker Training <education@docker.com>
```



ENV instruction

- Used to set environment variables in any container launched from the image
- Syntax

```
ENV <variable> <value>
```

Examples

```
ENV JAVA_HOME /usr/bin/java
```

```
ENV APP_PORT 8080
```



ADD instruction

- The ADD instruction copies new files or directories from a specified **source** and adds them to the container filesystem at a specified **destination**
- Syntax
ADD <src> <dest>
- The <src> path is relative to the directory container the Dockerfile
- If the <src> path is a directory, all files in the directory are copied. The directory itself is not copied
- You can specify multiple <src> directories

Example

```
ADD /src /myapp/src
```



COPY vs ADD

- Both instructions perform a near identical function
- ADD has the ability to auto unpack tar files
- ADD instruction also allows you to specify a URL for your content (although this is not recommended)
- Both instructions use a checksum against the files added. If the checksum is not equal then the test fails and the build cache will be invalidated
 - Because it means we have modified the files



Best practices for writing Dockerfiles

- Remember, each line in a Dockerfile creates a new layer
- You need to find the right balance between having lots of layers created for the image and readability of the Dockerfile
- Don't install unnecessary packages
- One ENTRYPOINT per Dockerfile
- Combine similar commands into one by using “&&” and “\”

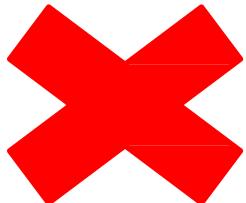
Example

```
RUN apt-get update && \
    apt-get install -y vim && \
    apt-get install -y curl
```



Best practices for writing Dockerfiles

- Use the caching system to your advantage
 - The order of statements is important
 - Add files that are least likely to change first and the ones most likely to change last
- The example below is not ideal because our build system does not know whether the requirements.txt file has changed



```
FROM ubuntu:14.04
MAINTAINER Docker Education Team <education@docker.com>
RUN apt-get update
RUN DEBIAN_FRONTEND=noninteractive apt-get install -y -q \
    python-all python-pip
ADD ./webapp /opt/webapp/
WORKDIR /opt/webapp
RUN pip install -qr requirements.txt
EXPOSE 5000
CMD ["python", "app.py"]
```



Best practices for writing Dockerfiles

- **Correct way** would be in the example below.
- `requirements.txt` is added in a separate step so Docker can cache more efficiently. If there's no change in the file the RUN pip install instruction does not have to execute and Docker can use the cache for that layer.
- The rest of the files are added afterwards



```
FROM ubuntu:14.04
MAINTAINER Docker Education Team <education@docker.com>
RUN apt-get update
RUN DEBIAN_FRONTEND=noninteractive apt-get install -y -q \
    python-all python-pip
ADD ./webapp/requirements.txt /tmp/requirements.txt
RUN pip install -qr /tmp/requirements.txt
ADD ./webapp /opt/webapp/
WORKDIR /opt/webapp
EXPOSE 5000
CMD ["python", "app.py"]
```



Module summary

- Images are made up of multiple layers
- Committing changes we make in a container as a new image is a simple way to create our own images but is not a very effective method as part of a development workflow
- A Dockerfile is the preferred method of creating images
- Key Dockerfile instructions we learnt about
 - RUN
 - CMD
 - ENTRYPOINT
 - COPY
- Key commands
 - docker build



Module 7: Volumes



Module objectives

In this module we will

- Explain what volumes are and what they are used for
- Learn the different methods of mounting a volume in a container
- Mount volumes during the `docker run` command and also in a Dockerfile
- Explain how data containers work
- Create some data containers



Volumes

*A **Volume** is a designated directory in a container, which is designed to persist data, independent of the container's life cycle*

- Volume changes are excluded when updating an image
- Persist when a container is deleted (kind of.. not documented dirs)
- Can be mapped to a host folder
- Can be shared between containers



Volumes and copy on write

- Volumes bypass the copy on write system
- Act as passthroughs to the host filesystem
- When you commit a container as a new image, the content of the volumes will not be brought into that image
- If a `RUN` instruction in a Dockerfile changes the content of a volume, those changes are not recorded either.



Uses of volumes

- De-couple the data that is stored, from the container which created the data
- Good for sharing data between containers
 - Can setup a data containers which has a volume you mount in other containers
 - Share directories between multiple containers
- Bypassing the copy on write system to achieve native disk I/O performance
- Share a host directory with a container
- Share a single file between the host and container



Mount a Volume

- Volumes can be mounted when running a container
- Use the `-v` option on `docker run`
- Volume paths specified must be absolute
- Can mount multiple volumes by using the `-v` option multiple times

Execute a new container and mount the folder `/myvolume` into its file system

```
docker run -d -P -v /myvolume nginx:1.7
```

Example of mounting multiple volumes

```
docker run -d -P -v /data/www -v /data/images nginx
```



Where are our volumes?

- Volumes exist independently from containers
- If a container is stopped, we can still access our volume
- To find where the volume is use `docker inspect` on the container

```
"Volumes": {  
    "/www/data": "/var/lib/docker/vfs/dir/1bfff820731cbbbaa2ec764269298deedd21f8bee55809cae309d03390892bf72a"  
},  
"VolumesRW": {  
    "/www/data": true  
}
```



Deleting a volume

- Volumes are not deleted when you delete a container
- To remove the volumes associated with a container use the `-v` option in the `docker rm` command

Delete a container and remove it's associated volumes

```
docker rm -v <container ID>
```



Deleting volumes

- If you created multiple containers which referenced the same volume, you will be able to access that volume as long as one of those containers is still present (running or stopped)
- When you remove the last container referencing a volume, that volume will be orphaned
- Orphaned volumes still exists but are very difficult to access and remove
- **Best practice:** When deleting a container, make sure you delete the volume associated with it, unless there are other containers using that volume



Mounting host folders to a volume

- When running a container, you can map folders on the host to a volume
- The files from the host folder will be present in the volume
- Changes made on the host are reflected inside the container volume
- **Syntax**
`docker run -v [host path]:[container path]:[rw|ro]`
- `rw` or `ro` controls the write status of the volume



Simple Example

- In the example below, files inside `/home/user/public_html` on the hosts will appear in the `/data/www` folder of the container
- If the host path or container path does not exist, it will be created
- If the container path is a folder with existing content, the files will be replaced by those from the host path

Mount the contents of the `public_html` folder on the hosts to the container volume at `/data/www`

```
docker run -d -v /home/user/public_html:/data/www ubuntu
```



Inspecting the mapped volume

- The volumes field from `docker inspect` will show the container volume being mapped to the host path specified during `docker run`

```
        "Paused": false,  
        "Pid": 0,  
        "Restarting": false,  
        "Running": false,  
        "StartedAt": "2015-05-20T13:56:48.403957851Z"  
    },  
    "Volumes": {  
        "/data/www": "/home/johnnytu/public_html"  
    },  
    "VolumesRW": {  
        "/data/www": false  
    }  
}
```



A more practical example with NGINX

- Let's run an NGINX container and have it serve web pages that we have on the host machine
- That way we can conveniently edit the page on the host instead of having to make changes inside the container
- Quick NGINX 101
 - NGINX starts with a one default server on port 80
 - Default location for pages is the `/usr/share/nginx/html` folder
 - By default the folder has an `index.html` file which is the welcome page



The NGINX welcome page

`index.html` page in
folder
`/usr/share/nginx/html`

Welcome to nginx!

If you see this page, the nginx web server is successfully installed and working. Further configuration is required.

For online documentation and support please refer to nginx.org.
Commercial support is available at nginx.com.

Thank you for using nginx.



Running the container

- Key aspects to remember when running the container
 - Use `-d` to run it in detached mode
 - Use `-P` to automatically map the container ports (more on this in Module 9)
- When we visit our server URL we should now see our `index.html` file inside our `public_html` folder instead of the default NGINX welcome page

Run an nginx container and map the our `public_html` folder to the volume at the `/usr/share/nginx/html` folder in the container. `<path>` is the path to `public_html`

```
docker run -d -P -v <path>:/usr/share/nginx/html nginx
```



Running the container (More Examples)

```
docker run -it -v ~/.bash_history:/.bash_history my image
```

```
docker run --rm -v $(pwd):/myfiles busybox sh
```

```
docker run --rm --volumes-from john1 -v $(pwd):/backup busybox tar cvf /  
backup/john2.tar /john
```



Use cases for mounting host directories

- You want to manage storage and snapshots yourself.
 - With LVM, or a SAN, or ZFS, or anything else!)
- You have a separate disk with better performance (SSD) or resiliency (EBS) than the system disk, and you want to put important data on that disk.
- You want to share your source directory between your host (where the source gets edited) and the container (where it is compiled or executed).
 - Good for testing purposes but not for production deployment



Volumes in Dockerfile

- VOLUME instruction creates a mount point
- Can specify arguments in a JSON array or string
- Cannot map volumes to host directories
- Volumes are initialized when the container is executed

String example

```
VOLUME /myvol
```

String example with multiple volumes

```
VOLUME /www/websitel.com /www/website2.com
```

JSON example

```
VOLUME ["myvol", "myvol2"]
```



Example Dockerfile with Volumes

- When we run a container from this image, the volume will be initialized along with any data in the specified location
- If we want to setup default files in the volume folder, the folder and file must be created first

```
FROM ubuntu:14.04

RUN apt-get update
RUN apt-get install -y vim \
    wget

RUN mkdir /data/myvol -p && \
    echo "hello world" > /data/myvol/testfile
VOLUME ["/data/myvol"]
```



Data containers

- A data container is a container created for the purpose of referencing one or many volumes
- Data containers don't run any application or process
- Used when you have persistent data that needs to be shared with other containers
- When creating a data container, you should give it a custom name to make it easier to reference



Custom container names

- By default, containers we create, have a randomly generated name
- To give your container a specific name, use the `--name` option on the `docker run` command
- Existing container can be renamed using the `docker rename` command
`docker rename <old name> <new name>`

Create a container and name it mynginx

```
docker run -d -P --name mynginx nginx
```

Rename the container called `happy_einstein` to `mycontainer`

```
docker rename happy_einstein mycontainer
```



Creating data containers

- We just need to run a container and specify a volume
- Should run a container using a lightweight image such as busybox
- No need to run any particular process, just run “true”

Run our data container using the busybox image

```
docker run --name mydata -v /data/app1 busybox true
```



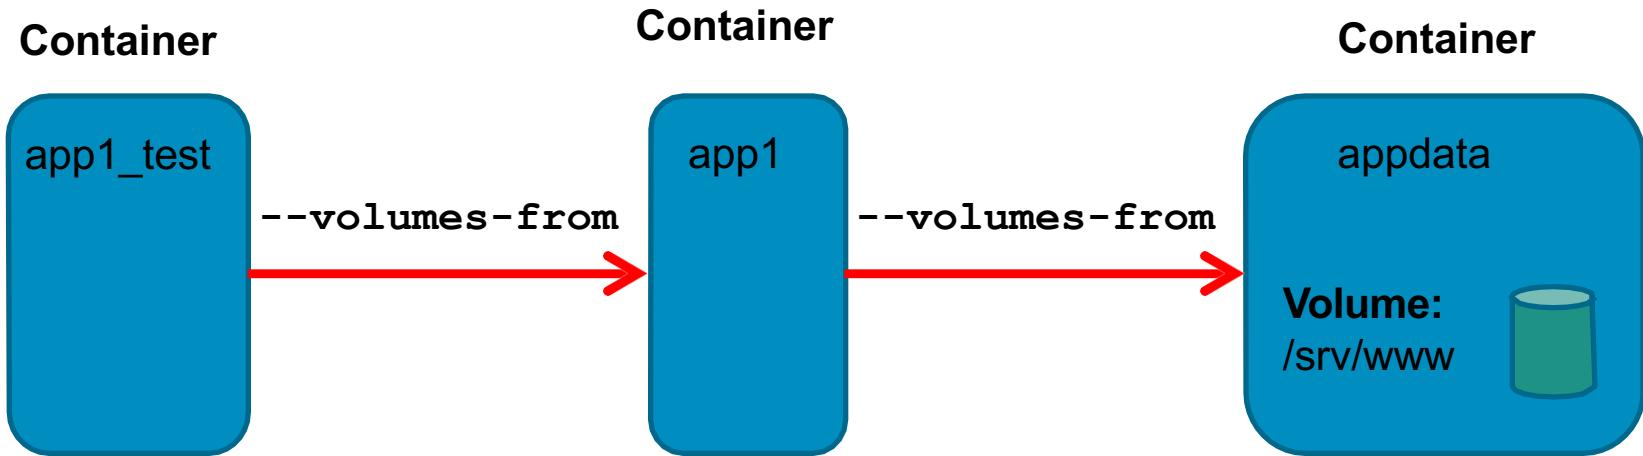
Using data containers

- Data containers can be used by other containers via the `--volumes-from` option in the `docker run` command
- Reference your data container by its container name. For example:
`--volumes-from datacontainer ...`

```
1 $ docker run --name appdata -v /data/app1 busybox
2 65d22a45d8ca11d77eed0faa8689d511a2265de1790e9bc41302378d0ac93c75
3
4 johnnytu@docker-ubuntu:~$ docker run -it --volumes-from appdata ubuntu:14.04
5 root@4ef756eeb298:#
6 root@4ef756eeb298:## cd /data/app1/
7 root@4ef756eeb298:/data/app1#
```



Chaining containers



- Container `app1` will mount the `/srv/www` volume from `appdata`
- Container `app1_test` will mount all volumes in `app1`, which will be the same `/srv/www` volume from `appdata`

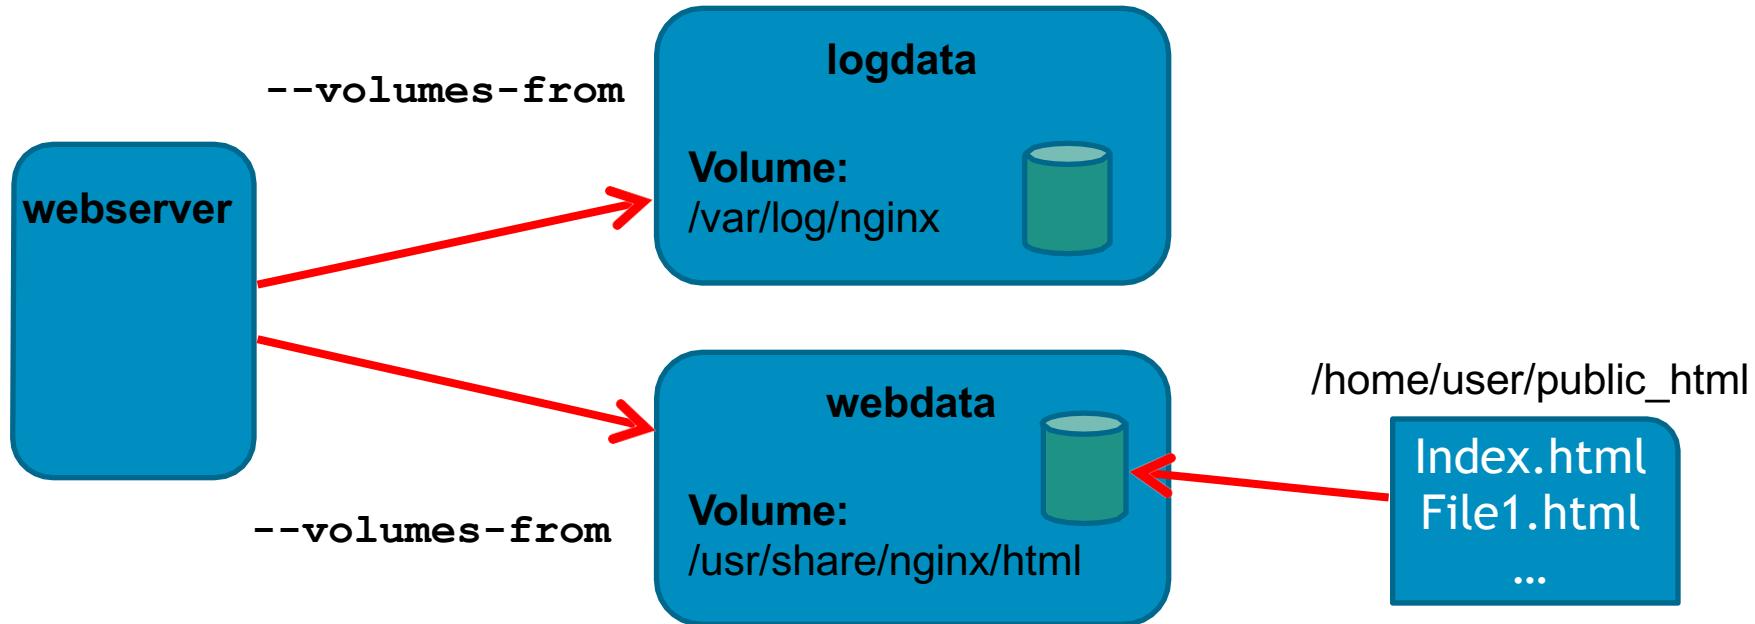


A more practical example

- Let's run an NGINX container but separate its operation across multiple containers
 - One container to handle the log files
 - One container to handle the web content that is to be served
 - One container to run NGINX



Diagram of example



Backup your data containers

- It's a good idea to back up data containers such as our logdata container, which has our NGINX log files
- Backups can be done with the following process:
 - Create a new container and mount the volumes from the data container
 - Mount a host directory as another volume on the container
 - Run the tar process to backup the data container volume onto your host folder

```
$ docker run --volumes-from logdata \
  -v /home/johnnytu/backups:/backup \
  ubuntu:14.04 \
  tar cvf /backup/nginxlogs.tar /var/log/nginx
```



Volumes defined in images

- Most images will have volumes defined in their Dockerfile
- Can check by using `docker inspect` command against the image
- `docker inspect` can be run against an image or a container
- To run against an image, specify either the image repository and tag or the image id.

Inspect the properties of the `ubuntu:14.04` image

```
docker inspect ubuntu:14.04
```

OR

```
docker inspect <image id>
```



Inspecting an image

```
  "Hostname": "d05bc6205be4",
  "Image": "e59ba510498bb53d2298ccc585b3140f4072f91070cd9b1c2bb504be87a4985b",
  "Labels": {},
  "MacAddress": "",
  "Memory": 0,
  "MemorySwap": 0,
  "NetworkDisabled": false,
  "OnBuild": [],
  "OpenStdin": false,
  "PortSpecs": null,
  "StdinOnce": false,
  "Tty": false,
  "User": "",
  "Volumes": {
    "/var/cache/nginx": {}
  },
  "WorkingDir": ""
},
"Created": "2015-04-30T23:05:04.539666944Z",
"DockerVersion": "1.6.0",
"Id": "42a3cf88f3f0cce2b4bfb2ed714eec5ee937525b4c7e0a0f70daff18c3f2ee92",
"Os": "linux",
"Parent": "c59ba510498bb53d2298ccc585b3140f4072f91070cd9b1c2bb504be87a4905b",
"Size": 0,
"VirtualSize": 132822837
}
]
```



Module summary

- Volumes can be mounted when we run a container during the `docker run` command or in a Dockerfile
- Volumes bypass the copy of write system
- We can map a host directory to a volume in a container
- A volume persists even after its container has been deleted
- A data container is a container created for the purpose of referencing one or many volumes



Module 8:

Container Networking



Module objectives

In this module we will

- Explain the Docker networking model for containers
- Learn how to map container ports to host ports manually and automatically
- Learn how to link containers together



Docker networking model

- Containers cannot have a public IPv4 address
- They are allocated a private address
- Services running on a container must be exposed port by port
- Container ports have to be mapped to the host port to avoid conflicts
- (Pre Libnetwork)



The docker0 bridge

- When Docker starts, it creates a virtual interface called `docker0` on the host machine
- `docker0` is assigned a random IP address and subnet from the private range defined by RFC 1918

```
johnnytu@docker-ubuntu:~$ ip a
...
...
3: docker0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc noqueue state
DOWN group default
    link/ether 56:84:7a:fe:97:99 brd ff:ff:ff:ff:ff:ff
        inet 172.17.42.1/16 scope global docker0
            valid_lft forever preferred_lft forever
        inet6 fe80::5484:7aff:fe:9799/64 scope link
            valid_lft forever preferred_lft forever
```



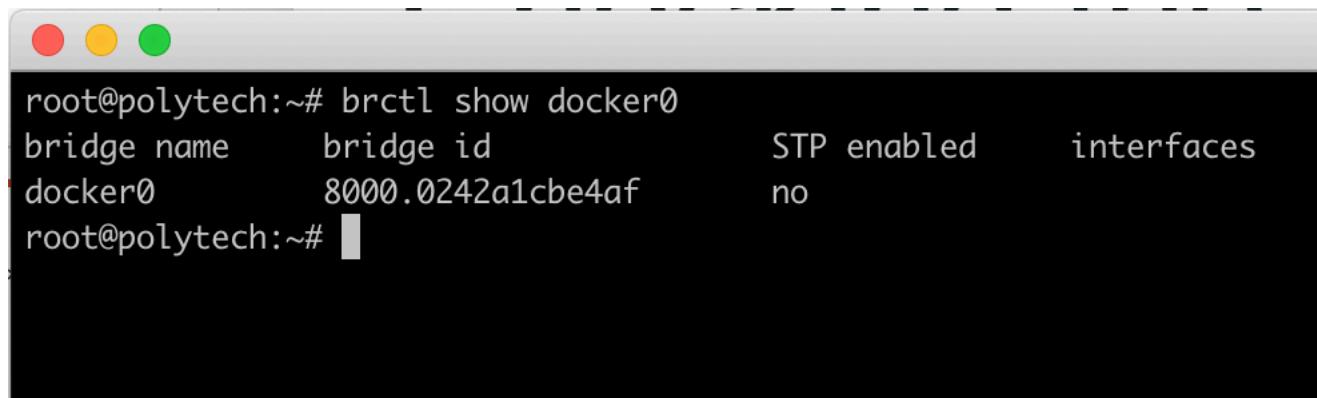
The docker0 bridge

- The `docker0` interface is a virtual Ethernet bridge interface
- It passes or switches packets between two connected devices just like a physical bridge or switch
 - Host to container
 - Container to container
- Each new container gets one interface that is automatically attached to the `docker0` bridge



Checking the bridge interface

- We can use the `brctl` (bridge control) command to check the interfaces on our `docker0` bridge
- **Install** `bridge-utils` package to get the command
`apt-get install bridge-utils`
- **Run**
`brctl show docker0`

A screenshot of a terminal window on a Mac OS X desktop. The window has a title bar with red, yellow, and green buttons. The main area of the terminal shows the output of the command `brctl show docker0`. The output is as follows:

```
root@polytech:~# brctl show docker0
bridge name      bridge id          STP enabled     interfaces
docker0          8000.0242a1cbe4af    no
root@polytech:~#
```



Checking the bridge interface

- Spin up some containers and then check the bridge again

```
root@polytech:~# docker run -it ubuntu:14.04
root@4044037139ee:/# root@polytech:~#
root@polytech:~#
root@polytech:~# docker run -it ubuntu:14.04
root@1f360ca99363:/#
root@1f360ca99363:/# root@polytech:~#
root@polytech:~# brctl show
bridge name      bridge id          STP enabled     interfaces
docker0          8000.0242a1cbe4af    no
                                         veth3ad11a9
                                         vethab15bc3
                                         vethd8b2a19
                                         vethdea8a8b
root@polytech:~#
```



Check container interface

- Get into the container terminal and run the `ip a` command

```
1. root@f8cba8135060: / (ssh)
root@polytech:~# ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP group default qlen 1000
    link/ether 08:00:27:8d:c0:4d brd ff:ff:ff:ff:ff:ff
    inet 10.0.2.15/24 brd 10.0.2.255 scope global eth0
        valid_lft forever preferred_lft forever
    inet6 fe80::a00:27ff:fe8d:c04d/64 scope link
        valid_lft forever preferred_lft forever
3: docker0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default
    link/ether 02:42:a1:cb:e4:af brd ff:ff:ff:ff:ff:ff
    inet 172.17.0.1/16 brd 172.17.255.255 scope global docker0
        valid_lft forever preferred_lft forever
    inet6 fe80::42:a1ff:fe:4af/64 scope link
        valid_lft forever preferred_lft forever
21: veth5ad11a9@ltz0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue master docker0 state UP group default
    link/ether f2:a7:dc:0b:51:5f brd ff:ff:ff:ff:ff:ff link-netnsid 0
    inet6 fe80::f0a7:dcff:fe0b:515f/64 scope link
        valid_lft forever preferred_lft forever
root@polytech:~#
```



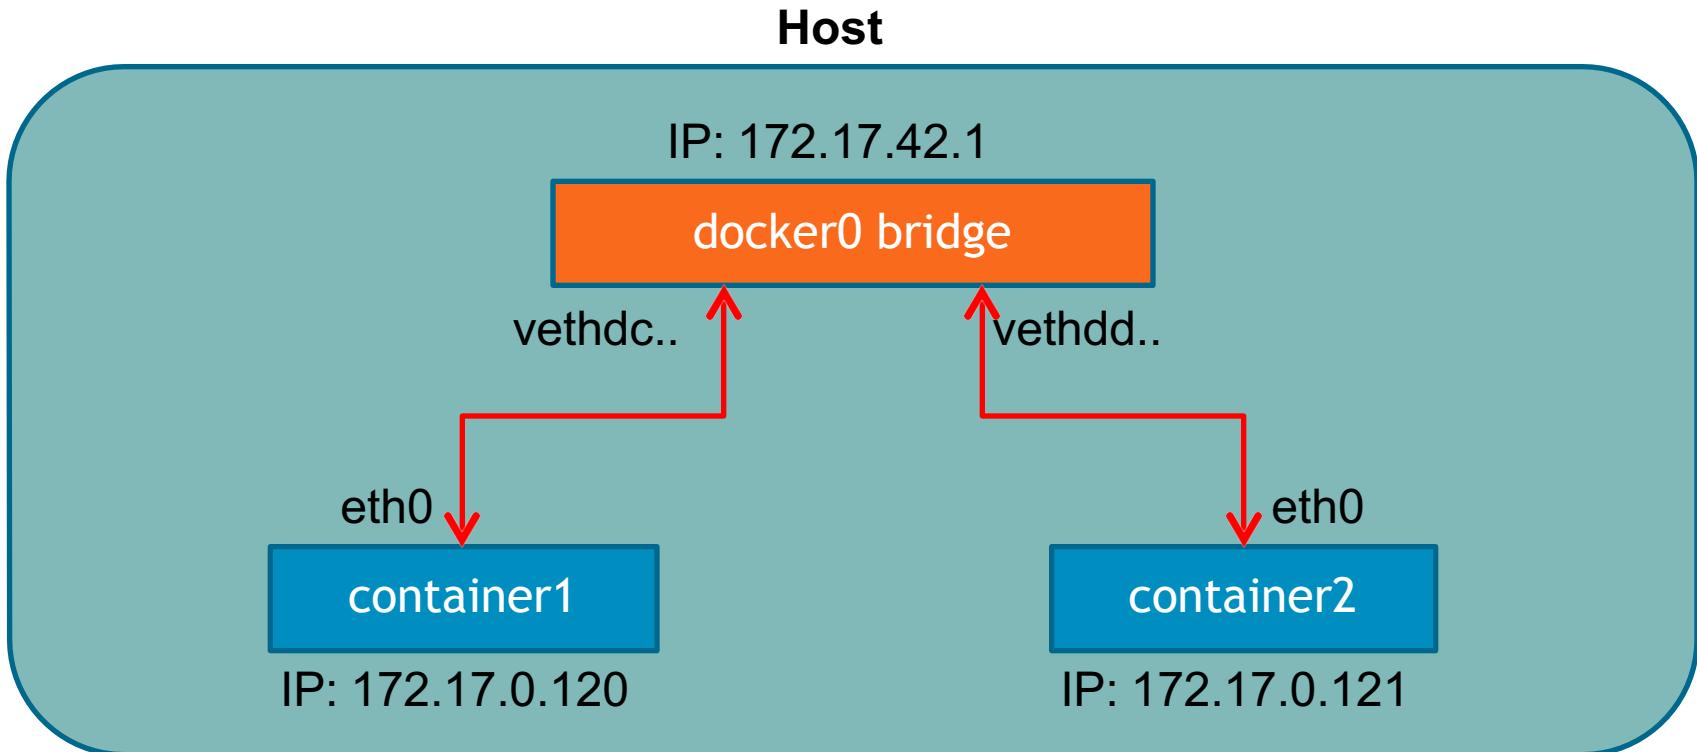
Check container networking properties

- Use `docker inspect` command and look for the `NetworkSettings` field

```
1. root@f4dbcfc:~# docker inspect --format='{{json .NetworkSettings}}' practical_merkle | jq
{
  "Bridge": "",
  "SandboxID": "9325594fb3ffc41e53a967c661b9e073eff30e1b2db990bd84e238dадec950da",
  "HairpinMode": false,
  "LinkLocalIPv6Address": "",
  "LinkLocalIPv6PrefixLen": 0,
  "Ports": {},
  "SandboxKey": "/var/run/docker/netns/9325594fb3ff",
  "SecondaryIPAddresses": null,
  "SecondaryIPv6Addresses": null,
  "EndpointID": "0a004d3a408d3b11beba406687d1cbd03abfa60a79c22cd7c0b3f9df279b5146",
  "Gateway": "172.17.0.1",
  "GlobalIPv6Address": "",
  "GlobalIPv6PrefixLen": 0,
  "IPAddress": "172.17.0.2",
  "IPPrefixLen": 16,
  "IPv6Gateway": "",
  "MacAddress": "02:42:ac:11:00:02",
  "Networks": {
    "bridge": {
      "IPAMConfig": null,
      "Links": null,
      "Aliases": null,
      "NetworkID": "9325594fb3ffc41e53a967c661b9e073eff30e1b2db990bd84e238dадec950da",
      "EndpointID": "0a004d3a408d3b11beba406687d1cbd03abfa60a79c22cd7c0b3f9df279b5146",
      "MacAddress": "02:42:ac:11:00:02",
      "IPv4Address": "172.17.0.2/16",
      "IPv6Address": null
    }
  }
}
```



Diagram of networking model



Mapping ports

- **Revision:** containers have their own network and IP address
- Map exposed container ports to ports on the host machine
- Ports can be manually mapped or auto mapped
- You can see the port mapping for each container on the docker ps output

```
root@polytech:~# docker ps
CONTAINER ID        IMAGE               COMMAND                  CREATED             STATUS              PORTS
2df7c7ba49e3        nginx              "nginx -g 'daemon of..."   6 seconds ago      Up 6 seconds      0.0.0.0:32769->80/tcp
7745040f1429        nginx              "nginx -g 'daemon of..."   17 seconds ago     Up 16 seconds     0.0.0.0:32768->80/tcp
f4dbcf617a7f        ubuntu:14.04      "/bin/bash"            14 minutes ago    Up 14 minutes
root@polytech:~#
```

1. root@f4dbcf617a7f: / (ssh)

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
2df7c7ba49e3	nginx	"nginx -g 'daemon of..."	6 seconds ago	Up 6 seconds	0.0.0.0:32769->80/tcp	elastic_knuth
7745040f1429	nginx	"nginx -g 'daemon of..."	17 seconds ago	Up 16 seconds	0.0.0.0:32768->80/tcp	vibrant_turing
f4dbcf617a7f	ubuntu:14.04	"/bin/bash"	14 minutes ago	Up 14 minutes		practical_merkle



Manual port mapping

- Uses the `-p` option (smaller case p) in the `docker run` command
- Syntax
`-p [host port]:[container port]`
- To map multiple ports, specify the `-p` option multiple times

Map port 8080 on the tomcat container to port 80 on the host

```
docker run -d -p 80:8080 tomcat
```

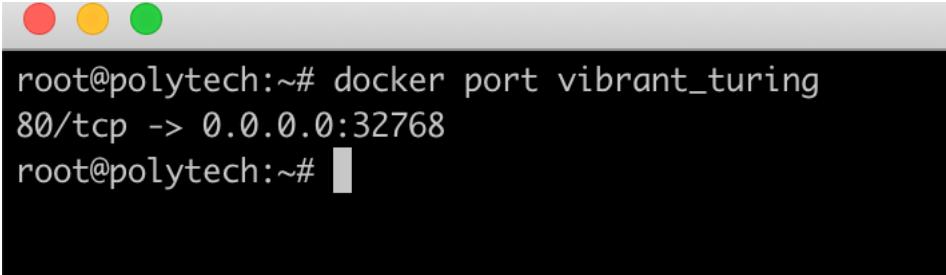
Map port on the host to port 80 on the nginx container and port 81 on the host to port 8080 on the nginx container

```
docker run -d -p 80:80 -p 81:8080 nginx
```



Docker port command

- docker ps output is not very ideal for display port mappings
- We can use the docker port command instead



```
root@polytech:~# docker port vibrant_turing
80/tcp -> 0.0.0.0:32768
root@polytech:~#
```



Automapping ports

- Use the `-P` option in `docker run` command
- Automatically maps exposed ports in the container to a port number in the host
- Host port numbers used go from 49153 to 65535
- Only works for ports defined in the Dockerfile `EXPOSE` instruction

Auto map ports exposed by the NGINX container to a port value on the host

```
docker run -d -P nginx:1.7
```



EXPOSE instruction

- Configures which ports a container will listen on at runtime
- Ports still need to be mapped when container is executed

```
FROM ubuntu:14.04
RUN apt-get update
RUN apt-get install -y nginx

EXPOSE 80 443

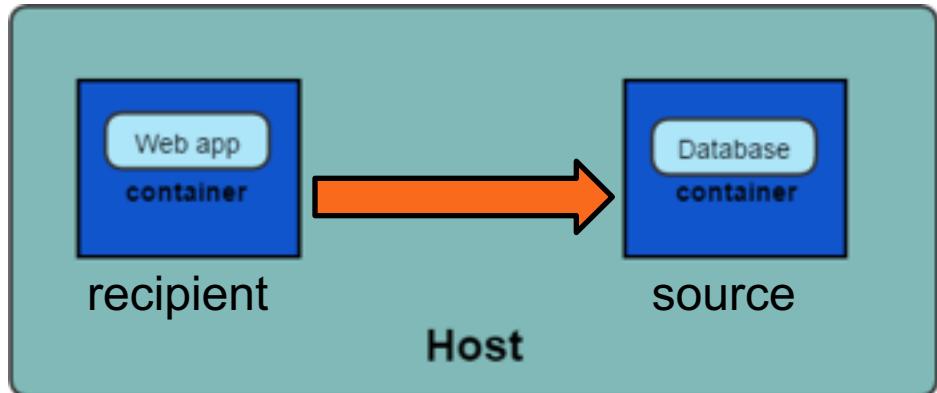
CMD ["nginx", "-g", "daemon off;"]
```



Linking Containers

Linking is a communication method between containers which allows them to securely transfer data from one to another

- Source and recipient containers
- Recipient containers have access to data on source containers
- Links are established based on container names



Uses of Linking

- Containers can talk to each other without having to expose ports to the host
- Essential for micro service application architecture
- Example:
 - Container with Tomcat running
 - Container with MySQL running
 - Application on Tomcat needs to connect to MySQL



Creating a Link

1. Create the source container first
 2. Create the recipient container and use the `--link` option
- **Best practice** – give your containers meaningful names
 - Format for linking
`name:alias`

Create the source container using the `postgres`

```
docker run -d --name database postgres
```

Create the recipient container and link it

```
docker run -d -P --name website --link database:db nginx
```



The underlying mechanism

- Linking provides a secure tunnel between the containers
- Docker will create a set of environment variables based on your --link parameter
- Docker also exposes the environment variables from the source container.
 - Only the variables created by Docker are exposed
 - Variables are prefixed by the link alias
 - ENV instruction in the container Dockerfile
 - Variables defined during docker run
- DNS lookup entry will be added to /etc/hosts file based on your alias



Environment variables in source container

- Let's inspect the variables defined in our "appserver" Tomcat container

```
$ docker exec appserver env
PATH=/usr/local/tomcat/bin:/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin
HOSTNAME=91794033c488
LANG=C.UTF-8
JAVA_VERSION=7u79
JAVA_DEBIAN_VERSION=7u79-2.5.5-1~deb8u1
CATALINA_HOME=/usr/local/tomcat
TOMCAT_MAJOR=8
TOMCAT_VERSION=8.0.22
TOMCAT_TGZ_URL=https://www.apache.org/dist/tomcat/tomcat-8/v8.0.22/bin/apache-tomcat-8.0.22.tar.gz
HOME=/root
```



Environment variables in recipient container

- Now we will check the environment variables in the Ubuntu container which is linked to our “appserver” container

```
$ docker exec client env
PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin
HOSTNAME=8a17932c348d
TERM=xterm
APP SERVER _PORT=tcp://172.17.0.133:8080
APP SERVER _PORT_8080_TCP=tcp://172.17.0.133:8080
APP SERVER _PORT_8080_TCP_ADDR=172.17.0.133
APP SERVER _PORT_8080_TCP_PORT=8080
APP SERVER _PORT_8080_TCP_PROTO=tcp
APP SERVER _NAME=/client/appserver
APP SERVER _ENV_LANG=C.UTF-8
APP SERVER _ENV_JAVA_VERSION=7u79
APP SERVER _ENV_JAVA_DEBIAN_VERSION=7u79-2.5.5-1~deb8u1
APP SERVER _ENV_CATALINA_HOME=/usr/local/tomcat
APP SERVER _ENV_TOMCAT_MAJOR=8
APP SERVER _ENV_TOMCAT_VERSION=8.0.22
APP SERVER _ENV_TOMCAT_TGZ_URL=https://www.apache.org/dist/tomcat/tomcat-8/v8.0.22/bin/apache-tomcat-8.0.22.tar.gz
HOME=/root
```

Connection information
created from the --link option

Created from source container



Module summary

- Docker containers run in a subnet provisioned by the docker0 bridge on the host machine
- Auto mapping of container ports to host ports only applies to the port numbers defined in the Dockerfile EXPOSE instruction
- Linking provides a secure method of communication between containers
- Key Dockerfile instructions
 - EXPOSE



Module summary

- Docker Compose makes it easier to manage micro service applications by making it easy to spin up and manage multiple containers
- Each service defined in the application is created as a container and can be scaled to multiple containers
- Docker Compose can create and run containers in a Swarm cluster



Module 9: Container Orchestration



Docker Swarm Kubernetes

- Go language
- Fork from Google Borg
- CNI Plugins
- Prototypage avec Minikube
- Offre SaaS : GKE, EKS, AKZ, ...



Further Information



Additional resources

- Docker homepage - <http://www.docker.com/>
- Docker Hub - <https://hub.docker.com>
- Docker documentation - <http://docs.docker.com/>
- Docker Getting Started Guide - <http://www.docker.com/gettingstarted/>
- Docker code on GitHub - <https://github.com/docker/docker>
- Kubernetes - <https://kubernetes.io/>
- Minikube - <https://kubernetes.io/docs/setup/minikube/>





THANK YOU